

```

/* controller/components/components.go */

// Package components defines the consolidated Entity-Component-System (ECS) components
// for the CPRA monitoring application. This design follows the principles of data-oriented
// design to maximize performance and minimize memory usage, as required for handling
// over one million concurrent monitors.
//
// By consolidating state, configuration, and jobs into a few coarse-grained components,
// we dramatically reduce the number of archetypes in the ECS world. This leads to:
//   - Improved cache locality and iteration speed.
//   - Reduced memory fragmentation.
//   - Simplified system logic by avoiding complex component additions/removals for state transitions.
//
// State management is handled via a bitfield in the MonitorState component, allowing for
// efficient, atomic updates to an entity's status without changing its archetype.
package components

import (
    "cpra/internal/jobs"
    "cpra/internal/loader/schema"
    "errors"
    "strings"
    "sync/atomic"
    "time"
)
)

// MonitorState consolidates all monitor state into a single component.
// This approach dramatically reduces archetype fragmentation and improves cache locality.
type MonitorState struct {
    // Entity identification
    Name string

    // State flags (bitfield for efficiency) - replaces multiple tag components
    Flags uint32

    // Timing data
    LastCheckTime    time.Time
    LastSuccessTime  time.Time
    NextCheckTime    time.Time

    // Error tracking
    ConsecutiveFailures int
    LastError          error

    // Pending action data
    PendingCode string
}

// State flag constants - replaces separate components like PulseNeeded, PulsePending, etc.
const (
    StateDisabled        uint32 = 1 << 0
    StatePulseNeeded     uint32 = 1 << 1
    StatePulsePending     uint32 = 1 << 2
    StatePulseFirstCheck  uint32 = 1 << 3
    StateInterventionNeeded uint32 = 1 << 4
    StateInterventionPending uint32 = 1 << 5
    StateCodeNeeded       uint32 = 1 << 6
    StateCodePending       uint32 = 1 << 7
    // Room for more states without adding components
)
)

// Efficient state management methods using atomic operations
func (m *MonitorState) IsDisabled() bool { return atomic.LoadUint32(&m.Flags)&StateDisabled != 0 }
func (m *MonitorState) IsPulseNeeded() bool {
    return atomic.LoadUint32(&m.Flags)&StatePulseNeeded != 0
}
func (m *MonitorState) IsPulsePending() bool {
    return atomic.LoadUint32(&m.Flags)&StatePulsePending != 0
}
func (m *MonitorState) IsPulseFirstCheck() bool {
    return atomic.LoadUint32(&m.Flags)&StatePulseFirstCheck != 0
}
func (m *MonitorState) IsInterventionNeeded() bool {
    return atomic.LoadUint32(&m.Flags)&StateInterventionNeeded != 0
}
func (m *MonitorState) IsInterventionPending() bool {
    return atomic.LoadUint32(&m.Flags)&StateInterventionPending != 0
}

```

```
func (m *MonitorState) IsCodeNeeded() bool {
    return atomic.LoadUint32(&m.Flags)&StateCodeNeeded != 0
}
func (m *MonitorState) IsCodePending() bool {
    return atomic.LoadUint32(&m.Flags)&StateCodePending != 0
}
func (m *MonitorState) SetDisabled(disabled bool) {
    if disabled {
        atomic.OrUint32(&m.Flags, StateDisabled)
    } else {
        atomic.AndUint32(&m.Flags, ^StateDisabled)
    }
}
func (m *MonitorState) SetPulseNeeded(needed bool) {
    if needed {
        atomic.OrUint32(&m.Flags, StatePulseNeeded)
    } else {
        atomic.AndUint32(&m.Flags, ^StatePulseNeeded)
    }
}
func (m *MonitorState) SetPulsePending(pending bool) {
    if pending {
        atomic.OrUint32(&m.Flags, StatePulsePending)
    } else {
        atomic.AndUint32(&m.Flags, ^StatePulsePending)
    }
}
func (m *MonitorState) SetPulseFirstCheck(firstCheck bool) {
    if firstCheck {
        atomic.OrUint32(&m.Flags, StatePulseFirstCheck)
    } else {
        atomic.AndUint32(&m.Flags, ^StatePulseFirstCheck)
    }
}
func (m *MonitorState) SetInterventionNeeded(needed bool) {
    if needed {
        atomic.OrUint32(&m.Flags, StateInterventionNeeded)
    } else {
        atomic.AndUint32(&m.Flags, ^StateInterventionNeeded)
    }
}
func (m *MonitorState) SetInterventionPending(pending bool) {
    if pending {
        atomic.OrUint32(&m.Flags, StateInterventionPending)
    } else {
        atomic.AndUint32(&m.Flags, ^StateInterventionPending)
    }
}
func (m *MonitorState) SetCodeNeeded(needed bool) {
    if needed {
        atomic.OrUint32(&m.Flags, StateCodeNeeded)
    } else {
        atomic.AndUint32(&m.Flags, ^StateCodeNeeded)
    }
}
func (m *MonitorState) SetCodePending(pending bool) {
    if pending {
        atomic.OrUint32(&m.Flags, StateCodePending)
    } else {
        atomic.AndUint32(&m.Flags, ^StateCodePending)
    }
}
// PulseConfig consolidates pulse configuration
type PulseConfig struct {
    Type          string
    Timeout       time.Duration
    Interval      time.Duration
    Retries       int
    MaxFailures  int
    Config        schema.PulseConfig
}
```

```

func (c *PulseConfig) Copy() *PulseConfig {
    if c == nil {
        return nil
    }
    cpy := &PulseConfig{
        Type:          strings.Clone(c.Type),
        Timeout:       c.Timeout,
        Interval:      c.Interval,
        Retries:       c.Retries,
        MaxFailures:   c.MaxFailures,
    }

    if c.Config != nil {
        cpy.Config = c.Config.Copy()
    }
    return cpy
}

// InterventionConfig consolidates intervention configuration
type InterventionConfig struct {
    Action      string
    MaxFailures int
    Target      schema.InterventionTarget
}

func (c *InterventionConfig) Copy() *InterventionConfig {
    if c == nil {
        return nil
    }
    cpy := &InterventionConfig{
        Action:          strings.Clone(c.Action),
        MaxFailures:    c.MaxFailures,
    }

    if c.Target != nil {
        cpy.Target = c.Target.Copy()
    }
    return cpy
}

// CodeConfig consolidates all code configurations instead of separate color components.
// This single component replaces RedCodeConfig, GreenCodeConfig, CyanCodeConfig, etc.
type CodeConfig struct {
    // Color-specific configurations stored as map instead of separate components
    Configs map[string]*ColorCodeConfig
}

type ColorCodeConfig struct {
    Dispatch    bool
    MaxFailures int
    Notify      string
    Config      schema.CodeNotification
}

func (c *ColorCodeConfig) Copy() *ColorCodeConfig {
    if c == nil {
        return nil
    }
    cpy := &ColorCodeConfig{
        Dispatch:    c.Dispatch,
        MaxFailures: c.MaxFailures,
        Notify:      strings.Clone(c.Notify),
    }

    if c.Config != nil {
        cpy.Config = c.Config.Copy()
    }
    return cpy
}

func (c *CodeConfig) Copy() *CodeConfig {
    if c == nil {
        return nil
    }
    cpy := &CodeConfig{
        Configs: make(map[string]*ColorCodeConfig),
    }

    for color, config := range c.Configs {
        cpy.Configs[color] = config.Copy()
    }
}

```

```

        return cpy
    }

// CodeStatus consolidates all code status instead of separate color status components
type CodeStatus struct {
    // Color-specific status stored as map
    Status map[string]*ColorCodeStatus
}

type ColorCodeStatus struct {
    LastStatus      string
    ConsecutiveFailures int
    LastAlertTime   time.Time
    LastSuccessTime time.Time
    LastError       error
}

func (s *ColorCodeStatus) SetSuccess(t time.Time) {
    s.LastStatus = "success"
    s.LastError = nil
    s.ConsecutiveFailures = 0
    s.LastSuccessTime = t
    s.LastAlertTime = t
}

func (s *ColorCodeStatus) SetFailure(err error) {
    s.LastStatus = "failed"
    s.LastError = err
    s.ConsecutiveFailures++
}

func (s *ColorCodeStatus) Copy() *ColorCodeStatus {
    if s == nil {
        return nil
    }
    cpy := &ColorCodeStatus{
        LastStatus:      strings.Clone(s.LastStatus),
        ConsecutiveFailures: s.ConsecutiveFailures,
        LastAlertTime:   s.LastAlertTime,
        LastSuccessTime: s.LastSuccessTime,
    }
    if s.LastError != nil {
        cpy.LastError = errors.New(s.LastError.Error())
    }
    return cpy
}

func (c *CodeStatus) Copy() *CodeStatus {
    if c == nil {
        return nil
    }
    cpy := &CodeStatus{
        Status: make(map[string]*ColorCodeStatus),
    }
    for color, status := range c.Status {
        cpy.Status[color] = status.Copy()
    }
    return cpy
}

// JobStorage consolidates all job storage instead of separate job components.
// This single component replaces PulseJob, InterventionJob, CodeJob, etc.
type JobStorage struct {
    PulseJob      jobs.Job
    InterventionJob jobs.Job
    CodeJobs      map[string]jobs.Job // Jobs for each code color
}

func (j *JobStorage) Copy() *JobStorage {
    if j == nil {
        return nil
    }
    cpy := &JobStorage{
        CodeJobs: make(map[string]jobs.Job),
    }
    if j.PulseJob != nil {
        cpy.PulseJob = j.PulseJob.Copy()
    }
    if j.InterventionJob != nil {
        cpy.InterventionJob = j.InterventionJob.Copy()
    }
}

```

```

        }
        for color, job := range j.CodeJobs {
            if job != nil {
                cpy.CodeJobs[color] = job.Copy()
            }
        }
        return cpy
    }

// Result components are used to convey job completion information back to the ECS.
// They are added to entities by the result handling logic and removed by the corresponding result
// system.

type PulseResult struct {
    Result jobs.Result
}

type InterventionResult struct {
    Result jobs.Result
}

type CodeResult struct {
    Result jobs.Result
}

/* controller/entities/mapper.go */

package entities

import (
    "cpra/internal/controller/components"
    "cpra/internal/jobs"
    "cpra/internal/loader/schema"
    "fmt"
    "strings"
    "time"
    "github.com/mlange-42/ark/ecs"
)

// EntityManager uses the new consolidated component design.
// This dramatically reduces the number of archetypes and improves performance.
type EntityManager struct {
    // Core consolidated components - only a few archetypes instead of dozens.
    MonitorState      *ecs.Map1[components.MonitorState]
    PulseConfig       *ecs.Map1[components.PulseConfig]
    InterventionConfig *ecs.Map1[components.InterventionConfig]
    CodeConfig        *ecs.Map1[components.CodeConfig]
    CodeStatus        *ecs.Map1[components.CodeStatus]
    JobStorage        *ecs.Map1[components.JobStorage]
}

// NewEntityManager creates a new consolidated entity manager.
func NewEntityManager(world *ecs.World) *EntityManager {
    return &EntityManager{
        MonitorState:      ecs.NewMap1[components.MonitorState](world),
        PulseConfig:       ecs.NewMap1[components.PulseConfig](world),
        InterventionConfig: ecs.NewMap1[components.InterventionConfig](world),
        CodeConfig:        ecs.NewMap1[components.CodeConfig](world),
        CodeStatus:        ecs.NewMap1[components.CodeStatus](world),
        JobStorage:        ecs.NewMap1[components.JobStorage](world),
    }
}

// CreateEntityFromMonitor creates an entity using the consolidated design.
func (e *EntityManager) CreateEntityFromMonitor(
    monitor *schema.Monitor,
    world *ecs.World) error {

    // Validation
    if world == nil {
        return fmt.Errorf("world cannot be nil")
    }
    if e == nil {
        return fmt.Errorf("EntityManager cannot be nil")
    }
    if monitor.Name == "" {
        fmt.Println(monitor, "name cannot be empty")
        return fmt.Errorf("monitor name cannot be empty")
    }
}

```

```

entity := world.NewEntity()
if !world.Alive(entity) {
    return fmt.Errorf("failed to create valid entity")
}

// Create consolidated MonitorState component
monitorState := &components.MonitorState{
    Name:          strings.Clone(monitor.Name),
    LastCheckTime: time.Now(),
    LastSuccessTime: time.Now(),
    NextCheckTime:  time.Now(),
}

// Set initial state flags
if monitor.Enabled {
    monitorState.SetPulseFirstCheck(true) // Equivalent to adding PulseFirstCheck
component
} else {
    monitorState.SetDisabled(true) // Equivalent to adding DisabledMonitor component
}

e.MonitorState.Add(entity, monitorState)

// Add pulse configuration
pulseConfig := &components.PulseConfig{
    Type:          strings.Clone(monitor.Pulse.Type),
    MaxFailures:  monitor.Pulse.MaxFailures,
    Timeout:      monitor.Pulse.Timeout,
    Interval:     monitor.Pulse.Interval,
    Config:       monitor.Pulse.Config.Copy(),
}
e.PulseConfig.Add(entity, pulseConfig)

// Create consolidated job storage
jobStorage := &components.JobStorage{
    CodeJobs: make(map[string]jobs.Job),
}

// Add pulse job
pulseJob, err := jobs.CreatePulseJob(monitor.Pulse, entity)
if err != nil {
    return err
}
jobStorage.PulseJob = pulseJob

// Add intervention if configured
if monitor.Intervention.Action != "" {
    maxFailures := 1
    if monitor.Intervention.MaxFailures > 0 {
        maxFailures = monitor.Intervention.MaxFailures
    }

    interventionConfig := &components.InterventionConfig{
        Action:      strings.Clone(monitor.Intervention.Action),
        Target:     monitor.Intervention.Target.Copy(),
        MaxFailures: maxFailures,
    }
    e.InterventionConfig.Add(entity, interventionConfig)

    // Add intervention job
    interventionJob, err := jobs.CreateInterventionJob(monitor.Intervention, entity)
    if err != nil {
        return err
    }
    jobStorage.InterventionJob = interventionJob
}

// Add consolidated code configuration instead of separate color components
if len(monitor.Codes) > 0 {
    codeConfig := &components.CodeConfig{
        Configs: make(map[string]*components.ColorCodeConfig),
    }
    codeStatus := &components.CodeStatus{
        Status: make(map[string]*components.ColorCodeStatus),
    }

    for color, config := range monitor.Codes {
        // Single consolidated entry instead of separate components

```

```

        codeConfig.Configs[color] = &components.ColorCodeConfig{
            Dispatch: config.Dispatch,
            Notify: strings.Clone(config.Notify),
            Config: config.Config.Copy(),
        }

        codeStatus.Status[color] = &components.ColorCodeStatus{
            LastAlertTime: time.Now(),
        }

        // Add code job to consolidated storage
        codeJob, err := jobs.CreateCodeJob(strings.Clone(monitor.Name), config,
entity, color)
        if err != nil {
            return err
        }
        jobStorage.CodeJobs[color] = codeJob
    }

    e.CodeConfig.Add(entity, codeConfig)
    e.CodeStatus.Add(entity, codeStatus)
}

e.JobStorage.Add(entity, jobStorage)

return nil
}

// EnableMonitor enables a monitor using consolidated state flags
func (e *EntityManager) EnableMonitor(entity ecs.Entity) {
    if state := e.MonitorState.Get(entity); state != nil {
        state.SetDisabled(false)
        state.SetPulseFirstCheck(true)
    }
}

// DisableMonitor disables a monitor using consolidated state flags
func (e *EntityManager) DisableMonitor(entity ecs.Entity) {
    if state := e.MonitorState.Get(entity); state != nil {
        state.SetDisabled(true)
        state.SetPulsePending(false)
        state.SetInterventionPending(false)
        state.SetCodePending(false)
    }
}

// GetMonitorState provides easy access to consolidated state
func (e *EntityManager) GetMonitorState(entity ecs.Entity) *components.MonitorState {
    return e.MonitorState.Get(entity)
}

/* controller/logger.go */

package controller

import (
    "context"
    "fmt"
    "log"
    "os"
    "strings"
    "time"
)
// LogLevel represents different logging levels
type LogLevel int

const (
    LogLevelDebug LogLevel = iota
    LogLevelInfo
    LogLevelWarn
    LogLevelError
    LogLevelFatal
)
var logLevelNames = map[LogLevel]string{
    LogLevelDebug: "DEBUG",
    LogLevelInfo: "INFO",
    LogLevelWarn: "WARN",
}

```

```

LogLevelError: "ERROR",
LogLevelFatal: "FATAL",
}

var LogLevelColors = map[LogLevel]string{
    LogLevelDebug: "\u001b[36m", // Cyan
    LogLevelInfo: "\u001b[32m", // Green
    LogLevelWarn: "\u001b[33m", // Yellow
    LogLevelError: "\u001b[31m", // Red
    LogLevelFatal: "\u001b[35m", // Magenta
}
const colorReset = "\u001b[0m"

// Logger provides structured logging with levels and context
type Logger struct {
    level      LogLevel
    component  string
    enableColor bool
    debugMode   bool
    prodMode    bool
    file        *os.File
    timezone    *time.Location
    tracer      *Tracer
}

// NewLogger creates a new logger instance
func NewLogger(component string, debugMode bool) *Logger {
    level := LogLevelInfo
    if debugMode {
        level = LogLevelDebug
    }

    // Check environment for production mode
    prodMode := strings.ToLower(os.Getenv("CPRA_ENV")) == "production"
    if prodMode {
        level = LogLevelWarn // More restrictive in production
    }

    // Enable colors for terminal output (disable in production)
    enableColor := !prodMode && isTerminal()

    // Get timezone from environment or use local timezone
    timezone := getTimezone()

    // Enable tracing in debug mode or if explicitly enabled
    enableTracing := debugMode || strings.ToLower(os.Getenv("CPRA_TRACING")) == "true"

    logger := &Logger{
        level:      level,
        component:  component,
        enableColor: enableColor,
        debugMode:   debugMode,
        prodMode:    prodMode,
        timezone:   timezone,
    }

    // Setup file logging for production
    if prodMode {
        logger.setupFileLogging()
    }

    // Setup tracing if enabled
    if enableTracing {
        logger.tracer = NewTracer(component, true)
    }

    return logger
}

// getTimezone returns the timezone to use for logging
func getTimezone() *time.Location {
    // Check environment variable first
    if tz := os.Getenv("CPRA_TIMEZONE"); tz != "" {
        if loc, err := time.LoadLocation(tz); err == nil {
            return loc
        }
        log.Printf("Warning: Invalid timezone '%s', using local timezone", tz)
    }
}

```

```

// Use local timezone as default
return time.Local
}

// isTerminal checks if we're running in a terminal
func isTerminal() bool {
    fileInfo, _ := os.Stdout.Stat()
    return (fileInfo.Mode() & os.ModeCharDevice) != 0
}

// setupFileLogging configures file output for production
func (l *Logger) setupFileLogging() {
    logFile := fmt.Sprintf("cptra-%s.log", time.Now().In(l.timezone).Format("2006-01-02"))
    file, err := os.OpenFile(logFile, os.O_CREATE|os.O_WRONLY|os.O_APPEND, 0644)
    if err != nil {
        log.Printf("Failed to open log file: %v", err)
        return
    }
    l.file = file
}

// Close closes the log file if open
func (l *Logger) Close() {
    if l.file != nil {
        l.file.Close()
    }
}

// formatMessage formats a log message with timestamp, level, and component
func (l *Logger) formatMessage(level LogLevel, msg string, args ...interface{}) string {
    // Use enhanced timestamp with timezone information - 12-hour format with AM/PM
    now := time.Now().In(l.timezone)
    timestamp := now.Format("2006-01-02 03:04:05.000 PM Z07:00") // 12-hour format with space and
AM/PM
    timezoneName := l.timezone.String()
    levelName := logLevelNames[level]

    formattedMsg := fmt.Sprintf(msg, args...)

    // Add tracing info if available
    traceInfo := ""
    if l.tracer != nil && l.tracer.enabled {
        stats := l.tracer.GetStats()
        if totalSpans, ok := stats["total_spans"].(int); ok && totalSpans > 0 {
            traceInfo = fmt.Sprintf(" [TRACE:spans=%d]", totalSpans)
        }
    }

    if l.enableColor {
        color := logLevelColors[level]
        return fmt.Sprintf("%s %s [%s%s%s] [%s]%s %s",
            timestamp, timezoneName, color, levelName, colorReset, l.component, traceInfo,
formattedMsg)
    }

    return fmt.Sprintf("%s %s [%s] [%s] %s %s",
        timestamp, timezoneName, levelName, l.component, traceInfo, formattedMsg)
}

// log writes a message at the specified level
func (l *Logger) log(level LogLevel, msg string, args ...interface{}) {
    if level < l.level {
        return
    }

    formatted := l.formatMessage(level, msg, args...)

    // Always output to stdout/stderr
    if level >= LogLevelError {
        fmt.Fprintf(os.Stderr, "%s\n", formatted)
    } else {
        fmt.Fprintf(os.Stdout, "%s\n", formatted)
    }

    // Also write to file in production
    if l.file != nil {
        fmt.Fprintf(l.file, "%s\n", formatted)
        l.file.Sync() // Ensure immediate write
    }
}

```

```

// Debug logs a debug message (only in debug mode)
func (l *Logger) Debug(msg string, args ...interface{}) {
    l.log(LogLevelDebug, msg, args...)
}

// Info logs an info message
func (l *Logger) Info(msg string, args ...interface{}) {
    l.log(LogLevelInfo, msg, args...)
}

// Warn logs a warning message
func (l *Logger) Warn(msg string, args ...interface{}) {
    l.log(LogLevelWarn, msg, args...)
}

// Error logs an error message
func (l *Logger) Error(msg string, args ...interface{}) {
    l.log(LogLevelError, msg, args...)
}

// Fatal logs a fatal message and exits
func (l *Logger) Fatal(msg string, args ...interface{}) {
    l.log(LogLevelFatal, msg, args...)
    os.Exit(1)
}

// WithContext creates a new logger with additional context
func (l *Logger) WithContext(context string) *Logger {
    return &Logger{
        level:      l.level,
        component:  fmt.Sprintf("%s:%s", l.component, context),
        enableColor: l.enableColor,
        debugMode:   l.debugMode,
        prodMode:    l.prodMode,
        file:       l.file,
        timezone:   l.timezone,
        tracer:     l.tracer,
    }
}

// StartTrace begins a new trace span with the logger's tracer
func (l *Logger) StartTrace(ctx context.Context, operation string) (context.Context, *TraceSpan) {
    if l.tracer != nil {
        return l.tracer.StartSpan(ctx, operation)
    }
    return ctx, nil
}

// FinishTrace completes a trace span
func (l *Logger) FinishTrace(span *TraceSpan, err error) {
    if l.tracer != nil {
        l.tracer.FinishSpan(span, err)
    }
}

// AddTraceTag adds a tag to a trace span
func (l *Logger) AddTraceTag(span *TraceSpan, key, value string) {
    if l.tracer != nil {
        l.tracer.AddSpanTag(span, key, value)
    }
}

// AddTraceMetadata adds metadata to a trace span
func (l *Logger) AddTraceMetadata(span *TraceSpan, key string, value interface{}) {
    if l.tracer != nil {
        l.tracer.AddSpanMetadata(span, key, value)
    }
}

// SetTraceEntity sets the entity ID for a trace span
func (l *Logger) SetTraceEntity(span *TraceSpan, entityID uint64) {
    if l.tracer != nil {
        l.tracer.SetSpanEntity(span, entityID)
    }
}

// GetTracingStats returns tracing statistics
func (l *Logger) GetTracingStats() map[string]interface{} {
    if l.tracer != nil {

```

```

        return l.tracer.GetStats()
    }
    return map[string]interface{}{"enabled": false}
}

// LogSystemPerformance logs system performance metrics
func (l *Logger) LogSystemPerformance(component string, duration time.Duration, entitiesProcessed int)
{
    if l.debugMode {
        rate := float64(entitiesProcessed) / duration.Seconds()
        l.Debug("Performance: %s processed %d entities in %v (%.1f/sec)",
            component, entitiesProcessed, duration, rate)
    }
}

// LogEntityOperation logs entity-level operations in debug mode
func (l *Logger) LogEntityOperation(operation string, entityID uint64, details string) {
    if l.debugMode {
        l.Debug("Entity[%d] %s: %s", entityID, operation, details)
    }
}

// LogWorkerPool logs worker pool statistics - debug only, completely silent otherwise
func (l *Logger) LogWorkerPool(poolName string, stats map[string]interface{}) {
    // Only log in debug mode, completely silent otherwise
}

// LogComponentState logs component state changes
func (l *Logger) LogComponentState(entityID uint32, component string, state string) {
    if l.debugMode {
        l.Debug("Entity[%d] %s -> %s", entityID, component, state)
    }
}

// LogChannelState logs channel buffer states
func (l *Logger) LogChannelState(channelName string, depth, capacity int) {
    if l.debugMode {
        utilization := float64(depth) / float64(capacity) * 100
        l.Debug("Channel[%s] depth: %d/%d (%.1f%% full)",
            channelName, depth, capacity, utilization)
    } else if depth == capacity {
        l.Warn("Channel[%s] is full (%d/%d)", channelName, depth, capacity)
    }
}

// LogJobExecution logs job execution details
func (l *Logger) LogJobExecution(jobType string, entityID uint64, duration time.Duration, success bool) {
    if l.debugMode {
        status := "SUCCESS"
        if !success {
            status = "FAILED"
        }
        l.Debug("Job[%s] Entity[%d] %s in %v", jobType, entityID, status, duration)
    }
}

// Global logger instances for different components
var (
    SystemLogger      *Logger
    SchedulerLogger   *Logger
    DispatchLogger    *Logger
    ResultLogger      *Logger
    WorkerPoolLogger  *Logger
    EntityLogger       *Logger
)

// InitializeLoggers sets up all component loggers
func InitializeLoggers(debugMode bool) {
    SystemLogger = NewLogger("SYSTEM", debugMode)
    SchedulerLogger = NewLogger("SCHEDULER", debugMode)
    DispatchLogger = NewLogger("DISPATCH", debugMode)
    ResultLogger = NewLogger("RESULT", debugMode)
    WorkerPoolLogger = NewLogger("WORKER", debugMode)
    EntityLogger = NewLogger("ENTITY", debugMode)
}

// CloseLoggers closes all logger files
func CloseLoggers() {
    loggers := []*Logger{

```

```

        SystemLogger, SchedulerLogger, DispatchLogger,
        ResultLogger, WorkerPoolLogger, EntityLogger,
    }

    for _, logger := range loggers {
        if logger != nil {
            logger.Close()
        }
    }
}

/* controller/memory.go */
package controller

import (
    "log"
    "runtime"
    "runtime/debug"
    "time"
)

// MemoryManager handles memory optimization and monitoring
type MemoryManager struct {
    maxMemory      uint64
    gcInterval     time.Duration
    lastGC         time.Time
    memoryStats    runtime.MemStats
    alertThreshold float64 // Percentage of max memory before alert
}

func NewMemoryManager(maxMemoryGB uint64, gcIntervalSeconds int) *MemoryManager {
    return &MemoryManager{
        maxMemory:      maxMemoryGB << 30, // Convert GB to bytes
        gcInterval:     time.Duration(gcIntervalSeconds) * time.Second,
        alertThreshold: 0.8, // Alert at 80% memory usage
    }
}

// MonitorMemory checks current memory usage and triggers cleanup if needed
func (m *MemoryManager) MonitorMemory() {
    runtime.ReadMemStats(&m.memoryStats)

    currentUsage := m.memoryStats.Alloc
    usagePercent := float64(currentUsage) / float64(m.maxMemory)

    if usagePercent > m.alertThreshold {
        log.Printf("HIGH MEMORY USAGE: %.2f%% (%d MB / %d MB)",
                  usagePercent*100,
                  currentUsage>>20,
                  m.maxMemory>>20)

        // Force garbage collection
        m.ForceGC()
    }

    // Periodic garbage collection
    if time.Since(m.lastGC) > m.gcInterval {
        runtime.GC()
        m.lastGC = time.Now()
    }
}

// ForceGC triggers immediate garbage collection with logging
func (m *MemoryManager) ForceGC() {
    before := m.memoryStats.Alloc
    runtime.GC()
    runtime.ReadMemStats(&m.memoryStats)
    after := m.memoryStats.Alloc

    freed := before - after
    log.Printf("Forced GC: freed %d MB (before: %d MB, after: %d MB)",
              freed>>20, before>>20, after>>20)

    m.lastGC = time.Now()
}

// GetMemoryStats returns current memory statistics
func (m *MemoryManager) GetMemoryStats() runtime.MemStats {
}

```

```

        runtime.ReadMemStats(&m.memoryStats)
        return m.memoryStats
    }

// SetMemoryLimit configures runtime memory limits
func (m *MemoryManager) SetMemoryLimit() {
    debug.SetMemoryLimit(int64(m.maxMemory))
    log.Printf("Memory limit set to: %d GB", m.maxMemory>>30)
}

// LogMemoryStats provides detailed memory information
func (m *MemoryManager) LogMemoryStats() {
    stats := m.GetMemoryStats()

    log.Printf("Memory Stats:")
    log.Printf(" Alloc: %d MB", stats.Alloc>>20)
    log.Printf(" TotalAlloc: %d MB", stats.TotalAlloc>>20)
    log.Printf(" Sys: %d MB", stats.Sys>>20)
    log.Printf(" NumGC: %d", stats.NumGC)
    log.Printf(" GCCPUFraction: %.4f", stats.GCCPUFraction)
}

/* controller/metrics.go */

package controller

import (
    "sync"
    "time"
)

// SystemMetrics holds performance metrics for a specific system
type SystemMetrics struct {
    SystemName      string
    TotalUpdates    int64
    TotalEntitiesProcessed int64
    TotalBatchesCreated int64
    TotalDuration   time.Duration
    MaxUpdateDuration time.Duration
    MinUpdateDuration time.Duration
    LastUpdateTime   time.Time
    StartTime       time.Time
}
}

// MetricsAggregator collects and aggregates metrics from all systems
type MetricsAggregator struct {
    mu      sync.RWMutex
    systems map[string]*SystemMetrics
}
}

// NewMetricsAggregator creates a new metrics aggregator
func NewMetricsAggregator() *MetricsAggregator {
    return &MetricsAggregator{
        systems: make(map[string]*SystemMetrics),
    }
}

// RegisterSystem registers a new system for metrics collection
func (ma *MetricsAggregator) RegisterSystem(systemName string) {
    ma.mu.Lock()
    defer ma.mu.Unlock()

    ma.systems[systemName] = &SystemMetrics{
        SystemName:      systemName,
        StartTime:       time.Now(),
        MinUpdateDuration: time.Hour, // Initialize to high value
    }
}

// RecordSystemUpdate records a system update with performance metrics
func (ma *MetricsAggregator) RecordSystemUpdate(systemName string, duration time.Duration,
entitiesProcessed int64, batchesCreated int64) {
    ma.mu.Lock()
    defer ma.mu.Unlock()

    metrics, exists := ma.systems[systemName]
    if !exists {
        // Auto-register system if not found
        metrics = &SystemMetrics{
            SystemName:      systemName,

```

```

        StartTime: time.Now(),
        MinUpdateDuration: time.Hour,
    }
    ma.systems[systemName] = metrics
}

metrics.TotalUpdates++
metrics.TotalEntitiesProcessed += entitiesProcessed
metrics.TotalBatchesCreated += batchesCreated
metrics.TotalDuration += duration
metrics.LastUpdateTime = time.Now()

// Update min/max durations
if duration > metrics.MaxUpdateDuration {
    metrics.MaxUpdateDuration = duration
}
if duration < metrics.MinUpdateDuration {
    metrics.MinUpdateDuration = duration
}
}

// GetSystemMetrics returns metrics for a specific system
func (ma *MetricsAggregator) GetSystemMetrics(systemName string) (*SystemMetrics, bool) {
    ma.mu.RLock()
    defer ma.mu.RUnlock()

    metrics, exists := ma.systems[systemName]
    if !exists {
        return nil, false
    }

    // Return a copy to avoid race conditions
    copy := *metrics
    return &copy, true
}

// GetAllMetrics returns metrics for all systems
func (ma *MetricsAggregator) GetAllMetrics() map[string]*SystemMetrics {
    ma.mu.RLock()
    defer ma.mu.RUnlock()

    result := make(map[string]*SystemMetrics)
    for name, metrics := range ma.systems {
        copy := *metrics
        result[name] = &copy
    }

    return result
}

// GetAggregateMetrics returns aggregate metrics across all systems
func (ma *MetricsAggregator) GetAggregateMetrics() AggregateMetrics {
    ma.mu.RLock()
    defer ma.mu.RUnlock()

    var aggregate AggregateMetrics
    aggregate.StartTime = time.Now()

    for _, metrics := range ma.systems {
        aggregate.TotalUpdates += metrics.TotalUpdates
        aggregate.TotalEntitiesProcessed += metrics.TotalEntitiesProcessed
        aggregate.TotalBatchesCreated += metrics.TotalBatchesCreated
        aggregate.TotalDuration += metrics.TotalDuration
        aggregate.SystemCount++

        if metrics.StartTime.Before(aggregate.StartTime) {
            aggregate.StartTime = metrics.StartTime
        }
        if metrics.MaxUpdateDuration > aggregate.MaxUpdateDuration {
            aggregate.MaxUpdateDuration = metrics.MaxUpdateDuration
        }
        if aggregate.MinUpdateDuration == 0 || (metrics.MinUpdateDuration <
aggregate.MinUpdateDuration && metrics.MinUpdateDuration > 0) {
            aggregate.MinUpdateDuration = metrics.MinUpdateDuration
        }
    }

    // Calculate averages
    if aggregate.TotalUpdates > 0 {
        aggregate.AvgUpdateDuration = aggregate.TotalDuration /

```

```

time.Duration(aggregate.TotalUpdates)
    aggregate.AvgEntitiesPerUpdate = float64(aggregate.TotalEntitiesProcessed) /
float64(aggregate.TotalUpdates)
    aggregate.AvgBatchesPerUpdate = float64(aggregate.TotalBatchesCreated) /
float64(aggregate.TotalUpdates)
}

// Calculate throughput
totalRuntime := time.Since(aggregate.StartTime)
if totalRuntime > 0 {
    aggregate.EntitiesPerSecond = float64(aggregate.TotalEntitiesProcessed) /
totalRuntime.Seconds()
    aggregate.UpdatesPerSecond = float64(aggregate.TotalUpdates) / totalRuntime.Seconds()
}

return aggregate
}

// AggregateMetrics holds aggregate performance metrics across all systems
type AggregateMetrics struct {
    SystemCount          int
    TotalUpdates         int64
    TotalEntitiesProcessed int64
    TotalBatchesCreated int64
    TotalDuration        time.Duration
    MaxUpdateDuration   time.Duration
    MinUpdateDuration   time.Duration
    AvgUpdateDuration   time.Duration
    AvgEntitiesPerUpdate float64
    AvgBatchesPerUpdate float64
    EntitiesPerSecond   float64
    UpdatesPerSecond    float64
    StartTime           time.Time
}
/* controller/optimized_controller.go */

package controller

import (
    "context"
    "cpra/internal/controller/systems"
    "cpra/internal/queue"
    "fmt"
    "log"
    "math"
    "os"
    "time"

    "cpra/internal/controller/entities"
    "cpra/internal/loader/streaming"
    "github.com/mlange-42/ark-tools/app"
    "github.com/mlange-42/ark/ecs"
)

// LoggerAdapter adapts the controller loggers to the systems interface.
type LoggerAdapter struct {
    logger interface {
        Info(format string, args ...interface{})
        Debug(format string, args ...interface{})
        Warn(format string, args ...interface{})
        Error(format string, args ...interface{})
        LogSystemPerformance(name string, duration time.Duration, count int)
    }
}

func (l *LoggerAdapter) Info(format string, args ...interface{}) { l.logger.Info(format, args...) }
func (l *LoggerAdapter) Debug(format string, args ...interface{}) { l.logger.Debug(format, args...) }
func (l *LoggerAdapter) Warn(format string, args ...interface{}) { l.logger.Warn(format, args...) }
func (l *LoggerAdapter) Error(format string, args ...interface{}) { l.logger.Error(format, args...) }
func (l *LoggerAdapter) LogSystemPerformance(name string, duration time.Duration, count int) {
    l.logger.LogSystemPerformance(name, duration, count)
}

func (l *LoggerAdapter) LogComponentState(entityID uint32, component string, action string) {
    l.logger.Debug("Entity[%d] component %s: %s", entityID, component, action)
}

// OptimizedController manages the ECS world and its systems using ark-tools.
type OptimizedController struct {
    app      *app.App
}

```

```

world *ecs.World
mapper *entities.EntityManager

pulseQueue queue.Queue
interventionQueue queue.Queue
codeQueue queue.Queue

pulsePool *queue.DynamicWorkerPool
interventionPool *queue.DynamicWorkerPool
codePool *queue.DynamicWorkerPool

// ECS Systems
pulseScheduleSystem *systems.BatchPulseScheduleSystem
pulseSystem *systems.BatchPulseSystem
pulseResultSystem *systems.BatchPulseResultSystem
interventionSystem *systems.BatchInterventionSystem
interventionResultSystem *systems.BatchInterventionResultSystem
codeSystem *systems.BatchCodeSystem
codeResultSystem *systems.BatchCodeResultSystem

config Config
running bool
}

// Config holds all configuration for the controller.
type Config struct {
    StreamingConfig streaming.StreamingConfig
    QueueCapacity uint64
    WorkerConfig queue.WorkerPoolConfig
    BatchSize int
    UpdateInterval time.Duration
}

// DefaultConfig returns a default configuration.
func DefaultConfig() Config {
    return Config{
        StreamingConfig: streaming.DefaultStreamingConfig(),
        QueueCapacity: 65536, // Must be a power of 2
        WorkerConfig: queue.DefaultWorkerPoolConfig(),
        BatchSize: 1000,
        UpdateInterval: 100 * time.Millisecond,
    }
}

// NewOptimizedController creates a new controller with the refactored systems using ark-tools.
func NewOptimizedController(config Config) *OptimizedController {
    // Create ark-tools app with initial capacity
    arkApp := app.New(1024)
    world := &arkApp.World
    mapper := entities.NewEntityManager(world)

    // Instantiate the new adaptive queue and dynamic worker pool.
    pulseQueue, err := queue.NewAdaptiveQueue(config.QueueCapacity)
    if err != nil {
        log.Fatalf("Failed to create pulse queue: %v", err)
    }
    interventionQueue, err := queue.NewAdaptiveQueue(config.QueueCapacity)
    if err != nil {
        log.Fatalf("Failed to create intervention queue: %v", err)
    }
    codeQueue, err := queue.NewAdaptiveQueue(config.QueueCapacity)
    if err != nil {
        log.Fatalf("Failed to create code queue: %v", err)
    }

    pulseLogger := log.New(os.Stdout, "[PulsePool] ", log.LstdFlags)
    pulsePool, err := queue.NewDynamicWorkerPool(pulseQueue, config.WorkerConfig, pulseLogger)
    if err != nil {
        log.Fatalf("Failed to create pulse worker pool: %v", err)
    }
    interventionLogger := log.New(os.Stdout, "[InterventionPool] ", log.LstdFlags)
    interventionPool, err := queue.NewDynamicWorkerPool(interventionQueue, config.WorkerConfig, interventionLogger)
    if err != nil {
        log.Fatalf("Failed to create intervention worker pool: %v", err)
    }
    codeLogger := log.New(os.Stdout, "[CodePool] ", log.LstdFlags)
    codePool, err := queue.NewDynamicWorkerPool(codeQueue, config.WorkerConfig, codeLogger)
    if err != nil {
        log.Fatalf("Failed to create code worker pool: %v", err)
    }
}

```

```

}

logger := &LoggerAdapter{logger: SystemLogger}

// Instantiate the refactored systems with dedicated queues and worker pools.
pulseRouter := pulsePool.GetRouter()
interventionRouter := interventionPool.GetRouter()
codeRouter := codePool.GetRouter()

pulseScheduleSystem := systems.NewBatchPulseScheduleSystem(world, logger)
pulseSystem := systems.NewBatchPulseSystem(world, pulseQueue, config.BatchSize, logger)
pulseResultSystem := systems.NewBatchPulseResultSystem(world, pulseRouter.PulseResultChan,
logger)

interventionSystem := systems.NewBatchInterventionSystem(world, interventionQueue,
config.BatchSize, logger)
interventionResultSystem := systems.NewBatchInterventionResultSystem(world,
interventionRouter.InterventionResultChan, logger)

codeSystem := systems.NewBatchCodeSystem(world, codeQueue, config.BatchSize, logger)
codeResultSystem := systems.NewBatchCodeResultSystem(world, codeRouter.CodeResultChan, logger)

arkApp.AddSystem(pulseScheduleSystem)
arkApp.AddSystem(pulseSystem)
arkApp.AddSystem(interventionSystem)
arkApp.AddSystem(codeSystem)
arkApp.AddSystem(pulseResultSystem)
arkApp.AddSystem(interventionResultSystem)
arkApp.AddSystem(codeResultSystem)

return &OptimizedController{
    app:                      arkApp,
    world:                     world,
    mapper:                    mapper,
    pulseQueue:                pulseQueue,
    interventionQueue:         interventionQueue,
    codeQueue:                 codeQueue,
    pulsePool:                 pulsePool,
    interventionPool:          interventionPool,
    codePool:                  codePool,
    pulseScheduleSystem:        pulseScheduleSystem,
    pulseSystem:                pulseSystem,
    pulseResultSystem:          pulseResultSystem,
    interventionSystem:         interventionSystem,
    interventionResultSystem:   interventionResultSystem,
    codeSystem:                 codeSystem,
    codeResultSystem:           codeResultSystem,
    config:                    config,
}
}

// LoadMonitors loads monitors using the streaming loader.
func (c *OptimizedController) LoadMonitors(ctx context.Context, filename string) error {
    loader := streaming.NewStreamingLoader(filename, c.world, c.config.StreamingConfig)
    stats, err := loader.Load(ctx)
    if err != nil {
        return fmt.Errorf("failed to load monitors: %w", err)
    }
    SystemLogger.Info("Successfully loaded %d monitors in %v (%.0f monitors/sec)",
        stats.TotalEntities, stats>LoadingTime, stats.CreationRate)
    if stats.PulseRate > 0 {
        limit := int(math.Ceil(stats.PulseRate * c.config.UpdateInterval.Seconds()))
        if limit < 1 {
            limit = 1
        }
        c.pulseSystem.SetMaxDispatch(limit)
    }
    return nil
}

// Start begins the main processing loop of the controller.
func (c *OptimizedController) Start(ctx context.Context) error {
    if c.running {
        return fmt.Errorf("controller already running")
    }
    c.pulsePool.Start()
    c.interventionPool.Start()
    c.codePool.Start()
    c.running = true
    go c.app.Run()
}

```

```

        SystemLogger.Info("Optimized controller started successfully")
        return nil
    }

// Stop gracefully shuts down the controller.
func (c *OptimizedController) Stop() {
    if !c.running {
        return
    }
    SystemLogger.Info("Stopping controller...")
    c.app.Finalize()
    c.running = false
    c.pulsePool.DrainAndStop()
    c.interventionPool.DrainAndStop()
    c.codePool.DrainAndStop()
    c.PrintShutdownMetrics()
    c.pulseQueue.Close()
    c.interventionQueue.Close()
    c.codeQueue.Close()
    SystemLogger.Info("Controller stopped")
}

// PrintShutdownMetrics logs queue, worker pool, and world statistics at shutdown.
func (c *OptimizedController) PrintShutdownMetrics() {
    logQueue := func(label string, stats queue.Stats) {
        SystemLogger.Info("%s Queue: depth=%d/%d enqueued=%d dequeued=%d dropped=%d", label,
        stats.QueueDepth, stats.Capacity, stats.Enqueue, stats.Dequeue, stats.Dropped)
        SystemLogger.Info("%s Queue timings: avg_wait=%v max_wait=%v window=%v", label,
        stats.AvgQueueTime, stats.MaxQueueTime, stats.SampleWindow)
        SystemLogger.Info("%s Queue rates: arrival=%.2f/s service=%.2f/s last_enqueue=%v
last_dequeue=%v", label, stats.EnqueueRate, stats.DequeueRate, stats.LastEnqueue, stats.LastDequeue)
    }
    logWorkers := func(label string, stats queue.WorkerPoolStats) {
        SystemLogger.Info("%s Workers: running=%d capacity=%d target=%d min=%d max=%d
waiting=%d", label, stats.RunningWorkers, stats.CurrentCapacity, stats.TargetWorkers,
stats.MinWorkers, stats.MaxWorkers, stats.WaitingTasks)
        SystemLogger.Info("%s Tasks: submitted=%d completed=%d pending_results=%d
scaling_events=%d last_scale=%v", label, stats.TasksSubmitted, stats.TasksCompleted,
stats.PendingResults, stats.ScalingEvents, stats.LastScaleTime)
    }
    SystemLogger.Info("== SHUTDOWN METRICS ==")

    logQueue("Pulse", c.pulseQueue.Stats())
    logQueue("Intervention", c.interventionQueue.Stats())
    logQueue("Code", c.codeQueue.Stats())

    logWorkers("Pulse", c.pulsePool.Stats())
    logWorkers("Intervention", c.interventionPool.Stats())
    logWorkers("Code", c.codePool.Stats())

    worldStats := c.world.Stats()
    SystemLogger.Info("World: entities_used=%d recycled=%d total=%d archetypes=%d components=%d
filters=%d locked=%t",
        worldStats.Entities.Used, worldStats.Entities.Recycled, worldStats.Entities.Total,
        len(worldStats.Archetypes), len(worldStats.ComponentTypes), worldStats.CachedFilters,
        worldStats.Locked)
    SystemLogger.Info("World memory: reserved=%dB used=%dB", worldStats.Memory,
    worldStats.MemoryUsed)
    SystemLogger.Info("====")
}

// GetWorld returns the ECS world for external access (e.g., testing, debugging).
func (c *OptimizedController) GetWorld() *ecs.World {
    return c.world
}

/* controller/recovery.go */

package controller

import (
    "cptra/internal/controller/entities"
    "log"
    "runtime/debug"
    "time"
)

```

```

// RecoverySystem provides system-level error recovery and health monitoring
type RecoverySystem struct {
    ErrorCount int
    LastError  time.Time
    MaxErrors  int
    ResetWindow time.Duration
    Mapper      *entities.EntityManager
}

func NewRecoverySystem(maxErrors int, resetWindow time.Duration) *RecoverySystem {
    return &RecoverySystem{
        MaxErrors:  maxErrors,
        ResetWindow: resetWindow,
    }
}

// SafeSystemUpdate wraps system updates with error recovery
func (r *RecoverySystem) SafeSystemUpdate(systemName string, updateFunc func() error) error {
    defer func() {
        if recovered := recover(); recovered != nil {
            r.ErrorCount++
            r.LastError = time.Now()

            log.Printf("PANIC in system %s: %v", systemName, recovered)
            log.Printf("Stack trace: %s", debug.Stack())

            // Circuit breaker logic
            if r.ErrorCount >= r.MaxErrors {
                log.Printf("System %s exceeded max errors (%d), entering degraded mode",
                          systemName, r.MaxErrors)
            }
        }
    }()

    // Reset error count if enough time has passed
    if time.Since(r.LastError) > r.ResetWindow {
        r.ErrorCount = 0
    }

    // Circuit breaker - prevent further damage if too many errors
    if r.ErrorCount >= r.MaxErrors {
        return nil // Skip execution
    }

    return updateFunc()
}

// ValidateEntityHealth checks entity component integrity
func (r *RecoverySystem) ValidateEntityHealth(w *ecs.World, entity ecs.Entity) bool {
    if !w.Alive(entity) {
        return false
    }

    // Check for required components
    if r.Mapper != nil {
        state := r.Mapper.GetMonitorState(entity)
        if state == nil {
            log.Printf("Entity %v missing MonitorState component", entity)
            return false
        }
        if state.Name == "" {
            log.Printf("Entity %v missing Name component", entity)
            return false
        }
    }
    return true
}

// CleanupOrphanedComponents removes components from dead entities
func (r *RecoverySystem) CleanupOrphanedComponents(w *ecs.World) {
    // This would need specific implementation based on component tracking
    // For now, log the cleanup intent
    log.Printf("Cleanup cycle: %d entities active", w.Stats().Entities.Used)
}

/* controller/systems/batch_code_result_system.go */

```

```

package systems

import (
    "cptra/internal/controller/components"
    "cptra/internal/jobs"
    "sync/atomic"
    "time"

    "github.com/mlange-42/ark/ecs"
)

// BatchCodeResultSystem processes the results of dispatched code alerts.
// It processes batches of results passed directly from the result router.
type BatchCodeResultSystem struct {
    world    *ecs.World
    logger   Logger

    // Mappers for efficient component access
    stateMapper *ecs.Map[components.MonitorState]
    ResultChan  <-chan []jobs.Result
}

// NewBatchCodeResultSystem creates a new BatchCodeResultSystem.
func NewBatchCodeResultSystem(world *ecs.World, results <-chan []jobs.Result, logger Logger) *BatchCodeResultSystem {
    return &BatchCodeResultSystem{
        world:    world,
        logger:   logger,
        stateMapper: ecs.NewMap[components.MonitorState](world),
        ResultChan: results,
    }
}

func (s *BatchCodeResultSystem) Initialize(w *ecs.World) {}

func (s *BatchCodeResultSystem) Update(w *ecs.World) {
    if s.ResultChan == nil {
        return
    }

    resultsBatches := make([][]jobs.Result, 0)
loop:
    for {
        select {
        case res, ok := <-s.ResultChan:
            if !ok {
                s.ResultChan = nil
                break loop
            }
            if len(res) == 0 {
                continue
            }
            resultsBatches = append(resultsBatches, res)
        default:
            break loop
        }
    }

    for _, res := range resultsBatches {
        s.ProcessBatch(res)
    }
}

// ProcessBatch processes a batch of code alert results.
func (s *BatchCodeResultSystem) ProcessBatch(results []jobs.Result) {
    startTime := time.Now()
    processedCount := 0

    for _, result := range results {
        ent := result.Entity()
        if !s.world.Alive(ent) {
            continue
        }

        state := s.stateMapper.Get(ent)
        if state == nil {
            continue
        }
    }
}

```

```

        // Ensure we are processing a pending code alert.
        if (atomic.LoadUint32(&state.Flags) & components.StateCodePending) == 0 {
            s.logger.Warn("Entity[%d] received CodeResult but was not in CodePending
state", ent.ID())
                continue
            }

            processedCount++

            // Extract color from the result payload.
            colorPayload, ok := result.Payload["color"]
            if !ok {
                s.logger.Warn("Entity[%d] has CodeResult with no color in payload", ent.ID())
                continue
            }
            color, ok := colorPayload.(string)
            if !ok {
                s.logger.Warn("Entity[%d] has CodeResult with invalid color payload type",
ent.ID())
                continue
            }

            if err := result.Error(); err != nil {
                s.logger.Error("Monitor '%s' %s alert failed to send: %v", state.Name, color,
err)
            } else {
                s.logger.Info("Monitor '%s' %s alert sent successfully.", state.Name, color)
            }

            // Unset the pending flag.
            atomic.AndUint32(&state.Flags, ^uint32(components.StateCodePending))
        }

        if processedCount > 0 {
            s.logger.LogSystemPerformance("BatchCodeResultSystem", time.Since(startTime),
processedCount)
        }
    }

// Finalize is a no-op for this system.
func (s *BatchCodeResultSystem) Finalize(w *ecs.World) {}

/* controller/systems/batch_code_system.go */

package systems

import (
    "cptra/internal/controller/components"
    "cptra/internal/jobs"
    "cptra/internal/queue"
    "reflect"
    "sync/atomic"
    "time"

    "github.com/mlange-42/ark/ecs"
)

// jobInfo is a helper struct to associate a job with its entity and color for batch processing.
type jobInfo struct {
    Entity ecs.Entity
    Job    jobs.Job
    Color  string
}

// BatchCodeSystem processes entities that need a code alert dispatched.
// It determines the correct color based on the entity's state and enqueues the job.
type BatchCodeSystem struct {
    world      *ecs.World
    queue     queue.Queue // Using a generic queue interface
    logger    Logger
    batchSize int

    // Filter for entities that require a code alert.
    filter     *ecs.Filter3[components.MonitorState, components.CodeConfig,
components.JobStorage]
    stateMapper *ecs.Map1[components.MonitorState]
}

```

```

// NewBatchCodeSystem creates a new BatchCodeSystem.
func NewBatchCodeSystem(world *ecs.World, q queue.Queue, batchSize int, logger Logger)
*BatchCodeSystem {
    return &BatchCodeSystem{
        world:      world,
        queue:      q,
        logger:     logger,
        batchSize:  batchSize,
        filter:     ecs.NewFilter3[components.MonitorState, components.CodeConfig,
components.JobStorage](world),
        stateMapper: ecs.NewMap1[components.MonitorState](world),
    }
}
func (s *BatchCodeSystem) Initialize(w *ecs.World) {
}

// Update finds and processes all monitors that need a code alert.
func (s *BatchCodeSystem) Update(w *ecs.World) {
    startTime := time.Now()
    stats := s.queue.Stats()
    if stats.Capacity > 0 && stats.QueueDepth >= int(float64(stats.Capacity)*0.9) {
        s.logger.Debug("Code queue saturated (%d/%d); deferring dispatch", stats.QueueDepth,
stats.Capacity)
    }

    query := s.filter.Query()

    free := stats.Capacity - stats.QueueDepth
    if free <= 0 {
        return
    }
    tokens := int(float64(free) * 0.8)
    if tokens <= 0 {
        tokens = free
    }
    if tokens <= 0 {
        tokens = 1
    }

    earlyExit := false

    jobsToProcess := make([]jobInfo, 0, s.batchSize)
    processedCount := 0

    for query.Next() {
        ent := query.Entity()
        state, _, jobStorage := query.Get()

        // Process only entities that need a code alert.
        if (atomic.LoadUint32(&state.Flags) & components.StateCodeNeeded) == 0 {
            continue
        }

        color := state.PendingCode
        if color == "" {
            // This should not happen if StateCodeNeeded is set, but as a safeguard:
            atomic.AndUint32(&state.Flags, ^uint32(components.StateCodeNeeded))
            continue
        }

        job, ok := jobStorage.CodeJobs[color]
        if !ok || isNilJob(job) {
            s.logger.Warn("Entity[%d] needs '%s' code alert, but no job is configured.", ent.ID(), color)
            // Clear the flag if no job is found to prevent spinning.
            atomic.AndUint32(&state.Flags, ^uint32(components.StateCodeNeeded))
            continue
        }

        jobsToProcess = append(jobsToProcess, jobInfo{Entity: ent, Job: job, Color: color})

        if len(jobsToProcess) >= tokens {
            s.processBatch(&jobsToProcess)
            processedCount += len(jobsToProcess)
            jobsToProcess = make([]jobInfo, 0, s.batchSize)
            earlyExit = true
            break
        }
    }
}

```

```

// Process any remaining entities
if earlyExit {
    query.Close()
}

if len(jobsToProcess) > 0 {
    s.processBatch(&jobsToProcess)
    processedCount += len(jobsToProcess)
}

if processedCount > 0 {
    s.logger.LogSystemPerformance("BatchCodeSystem", time.Since(startTime),
processedCount)
}

}

// processBatch attempts to enqueue a batch of jobs and updates entity states on success.
func (s *BatchCodeSystem) processBatch(jobsInfo []*jobInfo) {
    stats := s.queue.Stats()
    if stats.Capacity > 0 && stats.QueueDepth >= int(float64(stats.Capacity)*0.9) {
        s.logger.Debug("Code queue near capacity (%d/%d); skipping enqueue", stats.QueueDepth,
stats.Capacity)
        return
    }
    jobs := make([]interface{}, 0, len(*jobsInfo))
    submitted := make([]jobInfo, 0, len(*jobsInfo))
    for _, info := range *jobsInfo {
        if isNilJob(info.Job) {
            s.logger.Warn("Entity[%d] code job became nil before enqueue; skipping",
info.Entity.ID())
            continue
        }
        jobs = append(jobs, info.Job)
        submitted = append(submitted, info)
    }
    if len(jobs) == 0 {
        return
    }

    err := s.queue.EnqueueBatch(jobs)
    if err != nil {
        s.logger.Warn("Failed to enqueue code job batch, queue may be full: %v", err)
        return
    }

    for _, info := range submitted {
        if !s.world.Alive(info.Entity) {
            continue
        }
        state := s.stateMapper.Get(info.Entity)
        if state == nil {
            continue
        }

        for {
            flags := atomic.LoadUint32(&state.Flags)
            if flags&components.StateCodeNeeded == 0 {
                break
            }

            updated := (flags & ^uint32(components.StateCodeNeeded)) |
uint32(components.StateCodePending)
            if atomic.CompareAndSwapUint32(&state.Flags, flags, updated) {
                state.PendingCode = ""
                s.logger.Info("CODE DISPATCHED: %s (%s)", state.Name, info.Color)
                break
            }
        }
    }
}

// Finalize is a no-op for this system.
func (s *BatchCodeSystem) Finalize(w *ecs.World) {}

func isNilJob(job jobs.Job) bool {
    if job == nil {
        return true
    }
}

```

```

}
v := reflect.ValueOf(job)
switch v.Kind() {
    case reflect.Ptr, reflect.Interface, reflect.Slice, reflect.Map, reflect.Func, reflect.Chan:
        return v.IsNil()
    default:
        return false
}
}

/* controller/systems/batch_intervention_result_system.go */
package systems

import (
    "cpra/internal/controller/components"
    "cpra/internal/jobs"
    "sync/atomic"
    "time"

    "github.com/mlange-42/ark/ecs"
)

// BatchInterventionResultSystem processes completed intervention jobs.
// It processes batches of results passed directly from the result router.
type BatchInterventionResultSystem struct {
    world *ecs.World
    logger Logger

    // Mappers for efficient component access
    stateMapper *ecs.Map[components.MonitorState]
    ResultChan <-chan []jobs.Result
}

// NewBatchInterventionResultSystem creates a new BatchInterventionResultSystem.
func NewBatchInterventionResultSystem(world *ecs.World, results <-chan []jobs.Result, logger Logger) *BatchInterventionResultSystem {
    return &BatchInterventionResultSystem{
        world:      world,
        logger:     logger,
        stateMapper: ecs.NewMap[components.MonitorState](world),
        ResultChan: results,
    }
}

func (s *BatchInterventionResultSystem) Initialize(w *ecs.World) {}

func (s *BatchInterventionResultSystem) Update(w *ecs.World) {
    if s.ResultChan == nil {
        return
    }

    resultsBatches := make([][]jobs.Result, 0)
loop:
    for {
        select {
        case res, ok := <-s.ResultChan:
            if !ok {
                s.ResultChan = nil
                break loop
            }
            if len(res) == 0 {
                continue
            }
            resultsBatches = append(resultsBatches, res)
        default:
            break loop
        }
    }

    for _, res := range resultsBatches {
        s.ProcessBatch(res)
    }
}

// ProcessBatch processes a batch of intervention results.
func (s *BatchInterventionResultSystem) ProcessBatch(results []jobs.Result) {
    startTime := time.Now()
}

```

```

processedCount := 0

for _, result := range results {
    ent := result.Entity()
    if !s.world.Alive(ent) {
        continue
    }

    state := s.stateMapper.Get(ent)
    if state == nil {
        continue
    }

    // Ensure we are processing a pending intervention
    if (atomic.LoadUint32(&state.Flags) & components.StateInterventionPending) == 0 {
        s.logger.Warn("Entity[%d] received InterventionResult but was not in
InterventionPending state", ent.ID())
        continue
    }

    processedCount++
    state.LastCheckTime = time.Now()

    if result.Error() != nil {
        // --- FAILURE ---
        state.ConsecutiveFailures++
        state.LastError = result.Error()
        s.logger.Error("Monitor '%s' intervention failed: %v", state.Name,
state.LastError)
        s.triggerCode(ent, state, "red")
    } else {
        // --- SUCCESS ---
        s.logger.Info("Monitor '%s' intervention succeeded.", state.Name)
        state.ConsecutiveFailures = 0
        state.LastError = nil
        state.LastSuccessTime = state.LastCheckTime
        s.triggerCode(ent, state, "cyan")
    }

    // Unset the pending flag, regardless of outcome.
    atomic.AndUint32(&state.Flags, ^uint32(components.StateInterventionPending))
}

if processedCount > 0 {
    s.logger.LogSystemPerformance("BatchInterventionResultSystem", time.Since(startTime),
processedCount)
}
}

func (s *BatchInterventionResultSystem) triggerCode(entity ecs.Entity, state *components.MonitorState,
color string) {
    codeConfigMapper := ecs.NewMap[components.CodeConfig](s.world)
    if !codeConfigMapper.Has(entity) {
        return
    }
    codeConfig := codeConfigMapper.Get(entity)
    if _, ok := codeConfig.Configs[color]; ok {
        state.PendingCode = color
        atomic.OrUint32(&state.Flags, components.StateCodeNeeded)
        s.logger.Info("Monitor '%s' - flagging for %s alert code", state.Name, color)
    }
}

// Finalize is a no-op for this system.
func (s *BatchInterventionResultSystem) Finalize(w *ecs.World) {}

/* controller/systems/batch_intervention_system.go */

package systems

import (
    "cpra/internal/controller/components"
    "cpra/internal/queue"
    "sync/atomic"
    "time"
)

"github.com/mlange-42/ark/ecs"
)

```

```

// BatchInterventionSystem processes entities that need an intervention.
// It identifies entities with the StateInterventionNeeded flag, enqueues the corresponding job,
// and transitions the entity state to StateInterventionPending.
type BatchInterventionSystem struct {
    world      *ecs.World
    queue      queue.Queue // Using a generic queue interface
    logger     Logger
    batchSize  int

    // Filter for entities that require an intervention.
    filter      *ecs.Filter3[components.MonitorState, components.InterventionConfig,
components.JobStorage]
    monitorStateMapper *ecs.Map[components.MonitorState]
}

// NewBatchInterventionSystem creates a new BatchInterventionSystem.
func NewBatchInterventionSystem(world *ecs.World, q queue.Queue, batchSize int, logger Logger)
*BatchInterventionSystem {
    return &BatchInterventionSystem{
        world:          world,
        queue:         q,
        logger:        logger,
        batchSize:    batchSize,
        filter:        ecs.NewFilter3[components.MonitorState,
components.InterventionConfig, components.JobStorage](world),
        monitorStateMapper: ecs.NewMap[components.MonitorState](world),
    }
}

func (s *BatchInterventionSystem) Initialize(w *ecs.World) {}

// Update finds and processes all monitors that need an intervention.
func (s *BatchInterventionSystem) Update(w *ecs.World) {
    startTime := time.Now()
    stats := s.queue.Stats()
    if stats.Capacity > 0 && stats.QueueDepth >= int(float64(stats.Capacity)*0.9) {
        s.logger.Debug("Intervention queue saturated (%d/%d); deferring dispatch",
stats.QueueDepth, stats.Capacity)
    }

    query := s.filter.Query()

    free := stats.Capacity - stats.QueueDepth
    if free <= 0 {
        return
    }
    tokens := int(float64(free) * 0.8)
    if tokens <= 0 {
        tokens = free
    }
    if tokens <= 0 {
        tokens = 1
    }

    earlyExit := false

    jobsToQueue := make([]interface{}, 0, s.batchSize)
    entitiesToUpdate := make([]ecs.Entity, 0, s.batchSize)
    processedCount := 0

    for query.Next() {
        ent := query.Entity()
        state, _, jobStorage := query.Get()

        // Process only entities that need an intervention.
        if (atomic.LoadUint32(&state.Flags) & components.StateInterventionNeeded) == 0 {
            continue
        }

        if jobStorage.InterventionJob == nil {
            s.logger.Warn("Entity[%d] has InterventionNeeded state but no
InterventionJob", ent.ID())
            continue
        }

        jobsToQueue = append(jobsToQueue, jobStorage.InterventionJob)
        entitiesToUpdate = append(entitiesToUpdate, ent)

        if len(jobsToQueue) >= tokens {
            s.processBatch(&jobsToQueue, &entitiesToUpdate)
        }
    }
}

```

```

        processedCount += len(jobsToQueue)
        jobsToQueue = make([]interface{}, 0, s.batchSize)
        entitiesToUpdate = make([]ecs.Entity, 0, s.batchSize)
        earlyExit = true
        break
    }
}

// Process any remaining entities
if earlyExit {
    query.Close()
}

if len(jobsToQueue) > 0 {
    s.processBatch(&jobsToQueue, &entitiesToUpdate)
    processedCount += len(jobsToQueue)
}

if processedCount > 0 {
    s.logger.LogSystemPerformance("BatchInterventionSystem", time.Since(startTime),
processedCount)
}

}

// processBatch attempts to enqueue a batch of jobs and updates entity states on success.
func (s *BatchInterventionSystem) processBatch(jobs *[]interface{}, entities *[]ecs.Entity) {
    stats := s.queue.Stats()
    if stats.Capacity > 0 && stats.QueueDepth >= int(float64(stats.Capacity)*0.9) {
        s.logger.Debug("Intervention queue near capacity (%d/%d); skipping enqueue",
stats.QueueDepth, stats.Capacity)
        return
    }
    err := s.queue.EnqueueBatch(*jobs)
    if err != nil {
        s.logger.Warn("Failed to enqueue intervention job batch, queue may be full: %v", err)
        // Do not transition state if enqueue fails, allowing retry on the next tick.
        return
    }

    // If enqueue is successful, transition the state for all entities in the batch.
    for _, ent := range *entities {
        if !s.world.Alive(ent) {
            continue
        }
        state := s.monitorStateMapper.Get(ent)
        if state == nil {
            continue
        }

        for {
            flags := atomic.LoadUint32(&state.Flags)
            if flags&components.StateInterventionNeeded == 0 {
                break
            }

            updated := (flags & ^uint32(components.StateInterventionNeeded)) | uint32(components.StateInterventionPending)
            if atomic.CompareAndSwapUint32(&state.Flags, flags, updated) {
                s.logger.Info("INTERVENTION DISPATCHED: %s", state.Name)
                break
            }
        }
    }
}

// Finalize is a no-op for this system.
func (s *BatchInterventionSystem) Finalize(w *ecs.World) {}

/* controller/systems/batch_pulse_result_system.go */

package systems

import (
    "cpra/internal/controller/components"
    "cpra/internal/jobs"
    "sync/atomic"
    "time"
)

```

```

"github.com/mlange-42/ark/ecs"
)

// BatchPulseResultSystem processes completed pulse checks.
// It queries for entities with a PulseResult component and updates their state accordingly.
type BatchPulseResultSystem struct {
    world    *ecs.World
    logger   Logger

    // Mappers are used for efficient component access
    stateMapper      *ecs.Map1[components.MonitorState]
    configMapper     *ecs.Map1[components.PulseConfig]
    codeConfigMapper *ecs.Map1[components.CodeConfig]
    ResultChan       <-chan []jobs.Result
}

// NewBatchPulseResultSystem creates a new BatchPulseResultSystem.
func NewBatchPulseResultSystem(world *ecs.World, results <-chan []jobs.Result, logger Logger) *BatchPulseResultSystem {
    return &BatchPulseResultSystem{
        world:          world,
        logger:         logger,
        stateMapper:    ecs.NewMap1[components.MonitorState](world),
        configMapper:   ecs.NewMap1[components.PulseConfig](world),
        codeConfigMapper: ecs.NewMap1[components.CodeConfig](world),
        ResultChan:     results,
    }
}

func (s *BatchPulseResultSystem) Initialize(w *ecs.World) {
}

func (s *BatchPulseResultSystem) Update(w *ecs.World) {
    if s.ResultChan == nil {
        return
    }

    resultsBatches := make([][]jobs.Result, 0)
loop:
    for {
        select {
        case res, ok := <-s.ResultChan:
            if !ok {
                s.ResultChan = nil
                break loop
            }
            resultsBatches = append(resultsBatches, res)
        default:
            break loop
        }
    }

    for _, res := range resultsBatches {
        s.ProcessBatch(res)
    }
}

// ProcessBatch processes a batch of pulse results.
func (s *BatchPulseResultSystem) ProcessBatch(results []jobs.Result) {
    startTime := time.Now()
    processedCount := 0

    for _, result := range results {
        ent := result.Entity()
        if !s.world.Alive(ent) {
            continue
        }

        state := s.stateMapper.Get(ent)
        config := s.configMapper.Get(ent)

        if (atomic.LoadUint32(&state.Flags) & components.StatePulsePending) == 0 {
            s.logger.Warn("Entity[%d] received a PulseResult but was not in a PulsePending state.", ent.ID())
            continue
        }

        processedCount++
        state.LastCheckTime = time.Now()

        if result.Error() != nil {

```

```

// --- FAILURE ---
state.ConsecutiveFailures++
state.LastError = result.Error()
s.logger.Warn("Monitor '%s' pulse failed (%d/%d): %v", state.Name,
state.ConsecutiveFailures, config.MaxFailures, state.LastError)

        if state.ConsecutiveFailures >= config.MaxFailures {
            s.logger.Warn("Monitor '%s' reached max failures, triggering
intervention.", state.Name)
                atomic.OrUint32(&state.Flags, components.StateInterventionNeeded)
                state.ConsecutiveFailures = 0 // Reset after triggering
            } else if state.ConsecutiveFailures == 1 {
                s.triggerCode(ent, state, "yellow")
            }
        } else {
            // --- SUCCESS ---
            wasFailure := state.ConsecutiveFailures > 0
            if wasFailure {
                s.logger.Info("Monitor '%s' pulse recovered.", state.Name)
                s.triggerCode(ent, state, "green")
            }
            state.ConsecutiveFailures = 0
            state.LastError = nil
            state.LastSuccessTime = state.LastCheckTime
        }

        // Unset the pending flag, regardless of outcome.
atomic.AndUint32(&state.Flags, ^uint32(components.StatePulsePending))
}

if processedCount > 0 {
    s.logger.LogSystemPerformance("BatchPulseResultSystem", time.Since(startTime),
processedCount)
}
}

func (s *BatchPulseResultSystem) triggerCode(entity ecs.Entity, state *components.MonitorState, color
string) {
    codeConfig := s.codeConfigMapper.Get(entity)
    if codeConfig == nil {
        return
    }
    if _, ok := codeConfig.Configs[color]; ok {
        // TODO: This is a placeholder for a more robust CodeNeeded implementation
        // For now, we directly set the flag.
        state.PendingCode = color
        atomic.OrUint32(&state.Flags, components.StateCodeNeeded)
        s.logger.Info("Monitor '%s' - triggering %s alert code", state.Name, color)
    }
}

// Finalize is a no-op for this system.
func (s *BatchPulseResultSystem) Finalize(w *ecs.World) {}

/* controller/systems/batch_pulse_schedule_system.go */

package systems

import (
    "sync/atomic"
    "time"

    "cpra/internal/controller/components"
    "github.com/mlange-42/ark/ecs"
)

// BatchPulseScheduleSystem schedules pulse checks for entities that are due.
// It queries for monitors that are not disabled, not already pending a pulse check,
// and whose interval has passed since the last check.
// This system is a critical part of the monitoring pipeline, ensuring that checks
// are scheduled in a timely and efficient manner.
type BatchPulseScheduleSystem struct {
    world    *ecs.World
    logger  Logger

    // Filter for entities that are candidates for a pulse check.
    filter *ecs.Filter2[components.MonitorState, components.PulseConfig]
}

```

```

// NewBatchPulseScheduleSystem creates a new BatchPulseScheduleSystem.
func NewBatchPulseScheduleSystem(world *ecs.World, logger Logger) *BatchPulseScheduleSystem {
    return &BatchPulseScheduleSystem{
        world: world,
        logger: logger,
        filter: ecs.NewFilter2[components.MonitorState, components.PulseConfig](world),
    }
}

func (s *BatchPulseScheduleSystem) Initialize(w *ecs.World) {
}

// Update finds and schedules all monitors that are due for a pulse check.
func (s *BatchPulseScheduleSystem) Update(w *ecs.World) {
    start := time.Now()
    query := s.filter.Query()
    var scheduledCount int

    now := time.Now()

    for query.Next() {
        state, config := query.Get()

        for {
            flags := atomic.LoadUint32(&state.Flags)

            if (flags&components.StateDisabled != 0) || (flags&components.StatePulseNeeded != 0) || (flags&components.StatePulsePending != 0) {
                break
            }

            due := (flags&components.StatePulseFirstCheck != 0) || (now.Sub(state.LastCheckTime) >= config.Interval)
            if !due {
                break
            }

            updated := (flags | components.StatePulseNeeded) &^ components.StatePulseFirstCheck
            if atomic.CompareAndSwapUint32(&state.Flags, flags, updated) {
                scheduledCount++
                break
            }
        }
    }

    if scheduledCount > 0 {
        s.logger.LogSystemPerformance("BatchPulseScheduleSystem", time.Since(start), scheduledCount)
    }
}

// Finalize is a no-op for this system.
func (s *BatchPulseScheduleSystem) Finalize(w *ecs.World) {
    // Nothing to clean up
}

/* controller/systems/batch_pulse_system.go */

package systems

import (
    "cpra/internal/controller/components"
    "cpra/internal/queue"
    "sync/atomic"
    "time"

    "github.com/mlange-42/ark/ecs"
)

// BatchPulseSystem processes entities that need a pulse check.
// It identifies entities with the StatePulseNeeded flag, enqueues the corresponding job,
// and transitions the entity state to StatePulsePending.
type BatchPulseSystem struct {
    world      *ecs.World
    queue      queue.Queue // Using a generic queue interface
    logger     Logger
}

```

```

batchSize int
maxDispatch int

// Filter for entities that require a pulse check.
filter *ecs.Filter2[components.MonitorState, components.JobStorage]
monitorStateMapper *ecs.Map[components.MonitorState]
}

// NewBatchPulseSystem creates a new BatchPulseSystem.
func NewBatchPulseSystem(world *ecs.World, q queue.Queue, batchSize int, logger Logger)
*BatchPulseSystem {
    return &BatchPulseSystem{
        world:           world,
        queue:          q,
        logger:         logger,
        batchSize:      batchSize,
        filter:         ecs.NewFilter2[components.MonitorState, components.JobStorage]
(world),
        monitorStateMapper: ecs.NewMap[components.MonitorState](world),
    }
}

func (s *BatchPulseSystem) Initialize(w *ecs.World) {
}

func (s *BatchPulseSystem) SetMaxDispatch(n int) {
    s.maxDispatch = n
}

// Update finds and processes all monitors that need a pulse check.
func (s *BatchPulseSystem) Update(w *ecs.World) {
    startTime := time.Now()
    stats := s.queue.Stats()
    if stats.Capacity > 0 && stats.QueueDepth >= int(float64(stats.Capacity)*0.9) {
        s.logger.Debug("Pulse queue saturated (%d/%d); deferring dispatch", stats.QueueDepth,
stats.Capacity)
    }

    query := s.filter.Query()
    free := stats.Capacity - stats.QueueDepth
    if free <= 0 {
        return
    }
    tokens := int(float64(free) * 0.8)
    if tokens <= 0 {
        tokens = free
    }
    if tokens <= 0 {
        tokens = 1
    }
    if s.maxDispatch > 0 && tokens > s.maxDispatch {
        tokens = s.maxDispatch
    }

    earlyExit := false

    jobsToQueue := make([]interface{}, 0, s.batchSize)
    entitiesToUpdate := make([]ecs.Entity, 0, s.batchSize)
    processedCount := 0

    for query.Next() {
        ent := query.Entity()
        state, jobStorage := query.Get()

        // Process only entities that need a pulse check.
        if (atomic.LoadUint32(&state.Flags) & components.StatePulseNeeded) == 0 {
            continue
        }

        if jobStorage.PulseJob == nil {
            s.logger.Warn("Entity[%d] has PulseNeeded state but no PulseJob", ent.ID())
            continue
        }

        jobsToQueue = append(jobsToQueue, jobStorage.PulseJob)
        entitiesToUpdate = append(entitiesToUpdate, ent)

        if len(jobsToQueue) >= tokens {
            s.processBatch(&jobsToQueue, &entitiesToUpdate)
        }
    }
}

```

```

        processedCount += len(jobsToQueue)
        jobsToQueue = make([]interface{}, 0, s.batchSize)
        entitiesToUpdate = make([]ecs.Entity, 0, s.batchSize)
        earlyExit = true
        break
    }
}

// Process any remaining entities
if earlyExit {
    query.Close()
}

if len(jobsToQueue) > 0 {
    s.processBatch(&jobsToQueue, &entitiesToUpdate)
    processedCount += len(jobsToQueue)
}

if processedCount > 0 {
    s.logger.LogSystemPerformance("BatchPulseSystem", time.Since(startTime),
processedCount)
}

}

// processBatch attempts to enqueue a batch of jobs and updates entity states on success.
func (s *BatchPulseSystem) processBatch(jobs *[]interface{}, entities *[]ecs.Entity) {
    stats := s.queue.Stats()
    if stats.Capacity > 0 && stats.QueueDepth >= int(float64(stats.Capacity)*0.9) {
        s.logger.Debug("Pulse queue near capacity (%d/%d); skipping enqueue",
stats.QueueDepth, stats.Capacity)
        return
    }
    err := s.queue.EnqueueBatch(*jobs)
    if err != nil {
        s.logger.Warn("Failed to enqueue pulse job batch, queue may be full: %v", err)
        // Do not transition state if enqueue fails, allowing retry on the next tick.
        return
    }

    // If enqueue is successful, transition the state for all entities in the batch.
    now := time.Now()
    for _, ent := range *entities {
        if !s.world.Alive(ent) {
            continue
        }
        state := s.monitorStateMapper.Get(ent)
        if state == nil {
            continue
        }

        for {
            flags := atomic.LoadUint32(&state.Flags)
            if flags&components.StatePulseNeeded == 0 {
                break
            }

            updated := (flags & ^uint32(components.StatePulseNeeded)) |
uint32(components.StatePulsePending)
            if atomic.CompareAndSwapUint32(&state.Flags, flags, updated) {
                state.LastCheckTime = now
                break
            }
        }
    }
}

// Finalize is a no-op for this system.
func (s *BatchPulseSystem) Finalize(w *ecs.World) {}

/* controller/systems/errors.go */

package systems

import (
    "fmt"
    "time"
)

```

```

type ErrNoPulseJob struct {
    pulseType string
}

type ErrPulseJobTimeout struct {
    PulseType string
    Timeout    time.Duration
    Retries    int
    Err        error
}

func (e *ErrPulseJobTimeout) Error() string {
    return fmt.Sprintf("Job timeout for pulse type %s after %s (retried %d times): %s",
e.PulseType, e.Timeout, e.Retries, e.Err)
}

/* controller/systems/logger.go */

package systems

import "time"

// Logger interface for structured logging that all optimized systems can use
type Logger interface {
    Info(format string, args ...interface{})
    Debug(format string, args ...interface{})
    Warn(format string, args ...interface{})
    Error(format string, args ...interface{})
    LogSystemPerformance(name string, duration time.Duration, count int)
    LogComponentState(entityID uint32, component string, action string)
}
}

/* controller/systems/memory_efficient_system.go */

package systems

import (
    "fmt"
    "runtime"
    "time"

    "github.com/mlange-42/ark/ecs"
)

// MemoryEfficientSystem provides memory optimization utilities for ECS systems
type MemoryEfficientSystem struct {
    world          *ecs.World
    gcInterval     time.Duration
    lastGC         time.Time
    memoryThreshold int64

    // Memory statistics
    allocsBefore uint64
    allocsAfter  uint64
    gcCount       uint32
}

// MemoryConfig holds memory management configuration
type MemoryConfig struct {
    GCInterval     time.Duration // How often to check memory
    MemoryThreshold int64        // Memory threshold for forced GC (bytes)
    EnableProfiling bool         // Enable memory profiling
}

// MemoryStats holds memory statistics
type MemoryStats struct {
    Alloc          uint64        // Current allocated memory
    TotalAlloc     uint64        // Total allocated memory
    Sys            uint64        // System memory
    GCCount        uint32        // Number of GC runs
    LastGCTime    time.Time     // Last GC time
    GCPauseTotal   time.Duration // Total GC pause time
}

// NewMemoryEfficientSystem creates a new memory management system
func NewMemoryEfficientSystem(world *ecs.World, config MemoryConfig) *MemoryEfficientSystem {
    return &MemoryEfficientSystem{
        world:           world,
}

```

```

        gcInterval: config.GCInterval,
        memoryThreshold: config.MemoryThreshold,
        lastGC: time.Now(),
    }

// Update performs memory management tasks
func (mes *MemoryEfficientSystem) Update() {
    now := time.Now()

    // Check if it's time for memory management
    if now.Sub(mes.lastGC) < mes.gcInterval {
        return
    }

    // Get current memory stats
    var m runtime.MemStats
    runtime.ReadMemStats(&m)

    // Force GC if memory usage is high
    if int64(m.Alloc) > mes.memoryThreshold {
        mes.allocsBefore = m.Alloc
        runtime.GC()

        // Get stats after GC
        runtime.ReadMemStats(&m)
        mes.allocsAfter = m.Alloc
        mes.gcCount++

        fmt.Printf("Forced GC: Memory %d MB -> %d MB (freed %d MB)\n",
            mes.allocsBefore/1024/1024,
            m.Alloc/1024/1024,
            (mes.allocsBefore-m.Alloc)/1024/1024)
    }
}

mes.lastGC = now
}

// OptimizeEntityStorage optimizes entity storage by removing unused entities
func (mes *MemoryEfficientSystem) OptimizeEntityStorage() {
    // This would implement entity defragmentation if Ark supports it
    // For now, we just track the optimization request
    fmt.Println("Entity storage optimization requested (not implemented in Ark)")
}

// GetMemoryStats returns current memory statistics
func (mes *MemoryEfficientSystem) GetMemoryStats() MemoryStats {
    var m runtime.MemStats
    runtime.ReadMemStats(&m)

    return MemoryStats{
        Alloc: m.Alloc,
        TotalAlloc: m.TotalAlloc,
        Sys: m.Sys,
        GCCount: mes.gcCount,
        LastGCTime: mes.lastGC,
        GCPauseTotal: time.Duration(m.PauseTotalNs),
    }
}

// ForceGC forces garbage collection immediately
func (mes *MemoryEfficientSystem) ForceGC() {
    var m runtime.MemStats
    runtime.ReadMemStats(&m)
    mes.allocsBefore = m.Alloc

    runtime.GC()

    runtime.ReadMemStats(&m)
    mes.allocsAfter = m.Alloc
    mes.gcCount++
    mes.lastGC = time.Now()

    fmt.Printf("Manual GC: Memory %d MB -> %d MB (freed %d MB)\n",
        mes.allocsBefore/1024/1024,
        m.Alloc/1024/1024,
        (mes.allocsBefore-m.Alloc)/1024/1024)
}

// SetMemoryThreshold updates the memory threshold for automatic GC

```

```

func (mes *MemoryEfficientSystem) SetMemoryThreshold(threshold int64) {
    mes.memoryThreshold = threshold
}

// GetGCCount returns the number of forced GC runs
func (mes *MemoryEfficientSystem) GetGCCount() uint32 {
    return mes.gcCount
}

/* controller/systems/sys_test.go */

package systems

//
//import (
//    "cpra/internal/controller"
//    "cpra/internal/controller/components"
//    "cpra/internal/loader/loader"
//    "log"
//    "testing"
//)
//
//func TestPulseSystem_SchedulesJobs(t *testing.T) {
//    // Arrange
//    l := loader.NewLoader("yaml", "internal/loader/test.yaml.bak")
//    l.Load()
//    m := l.GetManifest()
//    w, err := controller.NewCPRaWorld(&m)
//    if err != nil {
//        log.Fatal(err)
//    }
//    entity := w.Mapper.CreateEntityFromMonitor(m)
//    w.AddComponent(entity, components.PulseConfig{ /* ... */ })
//    w.AddComponent(entity, components.Status{ /* ... */ })
//    // ... setup as needed
//
//    // Act
//    systems.PulseSystem{}.Update(w) // or .Step(w), etc.
//
//    // Assert
//    // Check if job scheduled, status updated, etc.
//    // Use t.Errorf or testify/assert
//}

```

```

/* controller/tracing.go */

package controller

import (
    "context"
    "fmt"
    "runtime"
    "sync"
    "time"

    "github.com/google/uuid"
)

// TraceSpan represents a single trace span
type TraceSpan struct {
    ID      string
    ParentID string
    Operation string
    StartTime time.Time
    EndTime   time.Time
    Duration  time.Duration
    Metadata map[string]interface{}
    Tags     map[string]string
    Success  bool
    Error    error
    Component string
    EntityID *uint64 // Optional entity ID for monitoring operations
}

// TraceContext holds trace information
type TraceContext struct {
    TraceID string
    SpanID  string
}

```

```

        Baggage map[string]string
    }

// Tracer manages trace collection and storage
type Tracer struct {
    mu      sync.RWMutex
    spans   map[string]*TraceSpan
    traces  map[string][]*TraceSpan // traceID -> spans
    enabled  bool
    component string
    logger   *Logger
}

// NewTracer creates a new tracer instance
func NewTracer(component string, enabled bool) *Tracer {
    // Create a simple logger without tracing to avoid circular dependency
    simpleLogger := &Logger{
        level:      LogLevelDebug,
        component:  fmt.Sprintf("TRACE:%s", component),
        enableColor: false,
        debugMode:   true,
        prodMode:    false,
        timezone:   time.Local,
        tracer:     nil, // No tracer to avoid recursion
    }

    return &Tracer{
        spans:     make(map[string]*TraceSpan),
        traces:    make(map[string][]*TraceSpan),
        enabled:   enabled,
        component: component,
        logger:    simpleLogger,
    }
}

// StartSpan begins a new trace span
func (t *Tracer) StartSpan(ctx context.Context, operation string) (context.Context, *TraceSpan) {
    if !t.enabled {
        return ctx, nil
    }

    var traceID, parentSpanID string

    // Check if there's an existing trace context
    if traceCtx, ok := ctx.Value("traceContext").(*TraceContext); ok {
        traceID = traceCtx.TraceID
        parentSpanID = traceCtx.SpanID
    } else {
        traceID = uuid.New().String()
    }

    spanID := uuid.New().String()

    span := &TraceSpan{
        ID:          spanID,
        ParentID:   parentSpanID,
        Operation:  operation,
        StartTime:  time.Now(),
        Metadata:   make(map[string]interface{}),
        Tags:        make(map[string]string),
        Component:  t.component,
    }

    // Add caller information
    if pc, file, line, ok := runtime.Caller(1); ok {
        span.Tags["caller.file"] = file
        span.Tags["caller.line"] = fmt.Sprintf("%d", line)
        span.Tags["caller.function"] = runtime.FuncForPC(pc).Name()
    }

    t.mu.Lock()
    t.spans[spanID] = span
    t.traces[traceID] = append(t.traces[traceID], span)
    t.mu.Unlock()

    // Create new context with trace information
    newTraceCtx := &TraceContext{
        TraceID: traceID,
        SpanID:  spanID,
        Baggage: make(map[string]string),
    }
}

```

```

}

newCtx := context.WithValue(ctx, "traceContext", newTraceCtx)
t.logger.Debug("Started span %s for operation %s (trace: %s, parent: %s)",
    spanID, operation, traceID, parentSpanID)

return newCtx, span
}

// FinishSpan completes a trace span
func (t *Tracer) FinishSpan(span *TraceSpan, err error) {
    if !t.enabled || span == nil {
        return
    }

    span.EndTime = time.Now()
    span.Duration = span.EndTime.Sub(span.StartTime)
    span.Success = err == nil
    span.Error = err

    t.mu.Lock()
    if existingSpan, exists := t.spans[span.ID]; exists {
        *existingSpan = *span
    }
    t.mu.Unlock()

    status := "SUCCESS"
    if err != nil {
        status = "ERROR"
    }

    t.logger.Debug("Finished span %s (%s) in %v [%s]",
        span.ID, span.Operation, span.Duration, status)

    if err != nil {
        t.logger.Debug("Span %s error: %v", span.ID, err)
    }
}

// AddSpanTag adds a tag to a span
func (t *Tracer) AddSpanTag(span *TraceSpan, key, value string) {
    if !t.enabled || span == nil {
        return
    }
    span.Tags[key] = value
}

// AddSpanMetadata adds metadata to a span
func (t *Tracer) AddSpanMetadata(span *TraceSpan, key string, value interface{}) {
    if !t.enabled || span == nil {
        return
    }
    span.Metadata[key] = value
}

// SetSpanEntity sets the entity ID for a span
func (t *Tracer) SetSpanEntity(span *TraceSpan, entityID uint64) {
    if !t.enabled || span == nil {
        return
    }
    span.EntityID = &entityID
    span.Tags["entity.id"] = fmt.Sprintf("%d", entityID)
}

// GetTrace returns all spans for a trace ID
func (t *Tracer) GetTrace(traceID string) []*TraceSpan {
    if !t.enabled {
        return nil
    }

    t.mu.RLock()
    defer t.mu.RUnlock()

    spans := make([]*TraceSpan, len(t.traces[traceID]))
    copy(spans, t.traces[traceID])
    return spans
}

// GetSpan returns a specific span by ID

```

```

func (t *Tracer) GetSpan(spanID string) *TraceSpan {
    if !t.enabled {
        return nil
    }

    t.mu.RLock()
    defer t.mu.RUnlock()

    if span, exists := t.spans[spanID]; exists {
        // Return a copy
        spanCopy := *span
        return &spanCopy
    }
    return nil
}

// GetStats returns tracing statistics
func (t *Tracer) GetStats() map[string]interface{} {
    if !t.enabled {
        return map[string]interface{}{"enabled": false}
    }

    t.mu.RLock()
    defer t.mu.RUnlock()

    totalSpans := len(t.spans)
    totalTraces := len(t.traces)

    var avgDuration time.Duration
    var successCount, errorCount int

    if totalSpans > 0 {
        var totalDuration time.Duration
        for _, span := range t.spans {
            if !span.EndTime.IsZero() {
                totalDuration += span.Duration
                if span.Success {
                    successCount++
                } else {
                    errorCount++
                }
            }
        }
        if totalSpans > 0 {
            avgDuration = totalDuration / time.Duration(totalSpans)
        }
    }

    return map[string]interface{}{
        "enabled": true,
        "total_spans": totalSpans,
        "total_traces": totalTraces,
        "success_count": successCount,
        "error_count": errorCount,
        "avg_duration": avgDuration.String(),
        "component": t.component,
    }
}

// Cleanup removes old spans to prevent memory leaks
func (t *Tracer) Cleanup(maxAge time.Duration) {
    if !t.enabled {
        return
    }

    t.mu.Lock()
    defer t.mu.Unlock()

    cutoff := time.Now().Add(-maxAge)
    var removedSpans, removedTraces int

    // Remove old spans
    for spanID, span := range t.spans {
        if span.EndTime.Before(cutoff) || (span.EndTime.IsZero() &&
span.StartTime.Before(cutoff)) {
            delete(t.spans, spanID)
            removedSpans++
        }
    }
}

```

```

// Remove empty traces
for traceID, spans := range t.traces {
    activeSpans := 0
    for _, span := range spans {
        if _, exists := t.spans[span.ID]; exists {
            activeSpans++
        }
    }
    if activeSpans == 0 {
        delete(t.traces, traceID)
        removedTraces++
    }
}

if removedSpans > 0 || removedTraces > 0 {
    t.logger.Debug("Cleaned up %d spans and %d traces", removedSpans, removedTraces)
}
}

// Global tracer instances
var (
    SystemTracer      *Tracer
    SchedulerTracer   *Tracer
    DispatchTracer    *Tracer
    ResultTracer      *Tracer
    WorkerPoolTracer  *Tracer
    EntityTracer      *Tracer
)
}

// InitializeTracers sets up all component tracers
func InitializeTracers(enabled bool) {
    SystemTracer = NewTracer("SYSTEM", enabled)
    SchedulerTracer = NewTracer("SCHEDULER", enabled)
    DispatchTracer = NewTracer("DISPATCH", enabled)
    ResultTracer = NewTracer("RESULT", enabled)
    WorkerPoolTracer = NewTracer("WORKER", enabled)
    EntityTracer = NewTracer("ENTITY", enabled)
}
}

// StartPeriodicCleanup starts a goroutine that periodically cleans up old traces
func StartPeriodicCleanup(interval, maxAge time.Duration) {
    tracers := []*Tracer{
        SystemTracer, SchedulerTracer, DispatchTracer,
        ResultTracer, WorkerPoolTracer, EntityTracer,
    }

    go func() {
        ticker := time.NewTicker(interval)
        defer ticker.Stop()

        for range ticker.C {
            for _, tracer := range tracers {
                if tracer != nil {
                    tracer.Cleanup(maxAge)
                }
            }
        }
    }()
}
}

/* jobs/jobs.go */

package jobs

import (
    "context"
    "cpra/internal/loader/schema"
    "fmt"
    "github.com/google/uuid"
    "github.com/mlange-42/ark/ecs"
    "github.com/moby/moby/api/types/container"
    "github.com/moby/moby/client"
    "net/http"
    "os"
    "strings"
    "time"
)
}

// Job defines the interface for any executable task in the system.
type Job interface {

```

```

Execute() Result
Copy() Job
GetEnqueueTime() time.Time
SetEnqueueTime(time.Time)
GetStartTime() time.Time
SetStartTime(time.Time)
}

// CreatePulseJob creates a new pulse job based on the provided schema.
func CreatePulseJob(pulseSchema schema.Pulse, jobID ecs.Entity) (Job, error) {
    timeout := pulseSchema.Timeout
    switch cfg := pulseSchema.Config.(type) {
    case *schema.PulseHTTPConfig:
        return &PulseHTTPJob{
            ID:        uuid.New(),
            Entity:   jobID,
            URL:      strings.Clone(cfg.Url),
            Method:   strings.Clone(cfg.Method),
            Timeout:  timeout,
            Retries:  cfg.Retries,
            Client:   http.Client{Timeout: timeout},
        }, nil
    case *schema.PulseTCPConfig:
        return &PulseTCPJob{
            ID:        uuid.New(),
            Entity:   jobID,
            Host:     strings.Clone(cfg.Host),
            Port:     cfg.Port,
            Timeout:  timeout,
            Retries:  cfg.Retries,
        }, nil
    case *schema.PulseICMPConfig:
        return &PulseICMPJob{
            ID:        uuid.New(),
            Entity:   jobID,
            Host:     strings.Clone(cfg.Host),
            Timeout:  timeout,
            Count:    cfg.Count,
        }, nil
    // ... other pulse job types
    default:
        return nil, fmt.Errorf("unknown pulse config type: %T for job creation",
pulseSchema.Config)
    }
}

// CreateInterventionJob creates a new intervention job based on the provided schema.
func CreateInterventionJob(interventionSchema schema.Intervention, jobID ecs.Entity) (Job, error) {
    retries := interventionSchema.Retries
    switch interventionSchema.Action {
    case "docker":
        return &InterventionDockerJob{
            ID:        uuid.New(),
            Entity:   jobID,
            Container: strings.Clone(interventionSchema.Target.
(*schema.InterventionTargetDocker).Container),
            Retries:  retries,
            Timeout:  interventionSchema.Target.
(*schema.InterventionTargetDocker).Timeout,
        }, nil
    default:
        return nil, fmt.Errorf("unknown intervention action : %T for job creation",
interventionSchema.Action)
    }
}

// CreateCodeJob creates a new code alert job based on the provided configuration.
func CreateCodeJob(monitor string, config schema.CodeConfig, jobID ecs.Entity, color string) (Job,
error) {
    // ... message creation logic ...
    message := "..."
    switch config.Notify {
    case "log":
        return &CodeLogJob{
            ID:        uuid.New(),
            File:     strings.Clone(config.Config.(*schema.CodeNotificationLog).File),
            Entity:   jobID,
            Monitor:  strings.Clone(monitor),
            Message:  message,
        }
    }
}

```

```

        Color: color,
    }, nil
    case "pagerduty":
        return &CodePagerDutyJob{
            ID:      uuid.New(),
            Entity: jobID,
            Monitor: strings.Clone(monitor),
            Message: message,
            Color: color,
        }, nil
    case "slack":
        return &CodeSlackJob{
            ID:      uuid.New(),
            Entity: jobID,
            Monitor: strings.Clone(monitor),
            Message: message,
            Color: color,
        }, nil
    case "email":
        return &CodeEmailJob{
            ID:      uuid.New(),
            Entity: jobID,
            Monitor: strings.Clone(monitor),
            Message: message,
            Color: color,
        }, nil
    case "webhook":
        return &CodeWebhookJob{
            ID:      uuid.New(),
            Entity: jobID,
            Monitor: strings.Clone(monitor),
            Message: message,
            Color: color,
        }, nil
    default:
        return nil, fmt.Errorf("unknown code notification type: %s for job creation",
config.Notify)
    }
}

// --- Pulse Job Implementations ---

type PulseHTTPJob struct {
    ID          uuid.UUID
    Entity      ecs.Entity
    URL         string
    Method      string
    Timeout     time.Duration
    Client      http.Client
    Retries     int
    EnqueueTime time.Time
    StartTime   time.Time
}

func (p *PulseHTTPJob) Execute() Result {
    var lastErr error
    attempts := p.Retries + 1
    payload := map[string]interface{}{"type": "pulse"}

    for i := 0; i < attempts; i++ {
        req, err := http.NewRequest(p.Method, p.URL, nil)
        if err != nil {
            return Result{ID: p.ID, Ent: p.Entity, Err: fmt.Errorf("failed to create http
request: %w", err), Payload: payload}
        }
        resp, err := p.Client.Do(req)
        if err != nil {
            lastErr = err
            time.Sleep(50 * time.Millisecond)
            continue
        }
        defer resp.Body.Close()
        if resp.StatusCode >= 200 && resp.StatusCode < 300 {
            return Result{ID: p.ID, Ent: p.Entity, Err: nil, Payload: payload}
        }
        lastErr = fmt.Errorf("received non-2xx status code: %s", resp.Status)
    }
    return Result{ID: p.ID, Ent: p.Entity, Err: fmt.Errorf("http check failed after %d attempt(s):
%w", attempts, lastErr), Payload: payload}
}

```

```

func (p *PulseHTTPJob) Copy() Job           { job := *p; return &job }
func (p *PulseHTTPJob) GetEnqueueTime() time.Time { return p.EnqueueTime }
func (p *PulseHTTPJob) SetEnqueueTime(t time.Time) { p.EnqueueTime = t }
func (p *PulseHTTPJob) GetStartTime() time.Time { return p.StartTime }
func (p *PulseHTTPJob) SetStartTime(t time.Time) { p.StartTime = t }

// PulseTCPJob is a placeholder for a TCP pulse job.
type PulseTCPJob struct {
    ID          uuid.UUID
    Entity      ecs.Entity
    Host        string
    Port        int
    Timeout     time.Duration
    Retries     int
    EnqueueTime time.Time
    StartTime   time.Time
}

func (p *PulseTCPJob) Execute() Result {
    // Mock implementation: does nothing and succeeds.
    return Result{ID: p.ID, Ent: p.Entity, Err: nil, Payload: map[string]interface{}{"type": "pulse", "driver": "tcp"}}
}

func (p *PulseTCPJob) Copy() Job           { job := *p; return &job }
func (p *PulseTCPJob) GetEnqueueTime() time.Time { return p.EnqueueTime }
func (p *PulseTCPJob) SetEnqueueTime(t time.Time) { p.EnqueueTime = t }
func (p *PulseTCPJob) GetStartTime() time.Time { return p.StartTime }
func (p *PulseTCPJob) SetStartTime(t time.Time) { p.StartTime = t }

// PulseICMPJob is a placeholder for an ICMP pulse job.
type PulseICMPJob struct {
    ID          uuid.UUID
    Entity      ecs.Entity
    Host        string
    Timeout     time.Duration
    Count       int
    EnqueueTime time.Time
    StartTime   time.Time
}

func (p *PulseICMPJob) Execute() Result {
    // Mock implementation: does nothing and succeeds.
    return Result{ID: p.ID, Ent: p.Entity, Err: nil, Payload: map[string]interface{}{"type": "pulse", "driver": "icmp"}}
}

func (p *PulseICMPJob) Copy() Job           { job := *p; return &job }
func (p *PulseICMPJob) GetEnqueueTime() time.Time { return p.EnqueueTime }
func (p *PulseICMPJob) SetEnqueueTime(t time.Time) { p.EnqueueTime = t }
func (p *PulseICMPJob) GetStartTime() time.Time { return p.StartTime }
func (p *PulseICMPJob) SetStartTime(t time.Time) { p.StartTime = t }

// --- Intervention Job Implementations ---

type InterventionDockerJob struct {
    ID          uuid.UUID
    Entity      ecs.Entity
    Container   string
    Timeout     time.Duration
    Retries     int
    EnqueueTime time.Time
    StartTime   time.Time
}

func (i *InterventionDockerJob) Execute() Result {
    payload := map[string]interface{}{"type": "intervention"}
    cli, err := client.NewClientWithOpts(client.FromEnv, client.WithAPIVersionNegotiation())
    if err != nil {
        return Result{ID: i.ID, Ent: i.Entity, Err: fmt.Errorf("failed to create docker client: %w", err), Payload: payload}
    }
    defer cli.Close()

    var lastErr error
    attempts := i.Retries + 1
    for attempt := 0; attempt < attempts; attempt++ {
        ctx, cancel := context.WithTimeout(context.Background(), i.Timeout)
        defer cancel()

```

```

        timeout := int(i.Timeout.Seconds())
        restartOptions := container.StopOptions{Timeout: &timeout}
        err := cli.ContainerRestart(ctx, i.Container, restartOptions)
        if err == nil {
            return Result{ID: i.ID, Ent: i.Entity, Err: nil, Payload: payload}
        }
        lastErr = err
    }
    return Result{ID: i.ID, Ent: i.Entity, Err: fmt.Errorf("docker intervention on '%s' failed
after %d attempt(s): %w", i.Container, attempts, lastErr), Payload: payload}
}

```

```

func (i *InterventionDockerJob) Copy() Job          { job := *i; return &job }
func (i *InterventionDockerJob) GetEnqueueTime() time.Time { return i.EnqueueTime }
func (i *InterventionDockerJob) SetEnqueueTime(t time.Time) { i.EnqueueTime = t }
func (i *InterventionDockerJob) GetStartTime() time.Time { return i.StartTime }
func (i *InterventionDockerJob) SetStartTime(t time.Time) { i.StartTime = t }

```

```
// --- Code Job Implementations ---
```

```

type CodeLogJob struct {
    ID      uuid.UUID
    Entity   ecs.Entity
    File     string
    Message  string
    Monitor  string
    Color    string
    EnqueueTime time.Time
    StartTime time.Time
}

```

```

func (c *CodeLogJob) Execute() Result {
    payload := map[string]interface{}{"type": "code", "color": c.Color}
    f, err := os.OpenFile(c.File, os.O_APPEND|os.O_CREATE|os.O_WRONLY, 0644)
    if err != nil {
        return Result{ID: c.ID, Ent: c.Entity, Err: err, Payload: payload}
    }
    defer f.Close()

    timestamp := time.Now().Format("2006-01-02 15:04:05.000 Z07:00")
    logLine := fmt.Sprintf("%s [%s] %s\n", timestamp, c.Monitor, c.Message)

    _, err = f.WriteString(logLine)
    return Result{ID: c.ID, Ent: c.Entity, Err: err, Payload: payload}
}

```

```

func (c *CodeLogJob) Copy() Job          { job := *c; return &job }
func (c *CodeLogJob) GetEnqueueTime() time.Time { return c.EnqueueTime }
func (c *CodeLogJob) SetEnqueueTime(t time.Time) { c.EnqueueTime = t }
func (c *CodeLogJob) GetStartTime() time.Time { return c.StartTime }
func (c *CodeLogJob) SetStartTime(t time.Time) { c.StartTime = t }

```

```
// CodePagerDutyJob is a placeholder for a PagerDuty notification job.
```

```

type CodePagerDutyJob struct {
    ID      uuid.UUID
    Entity   ecs.Entity
    Monitor  string
    Message  string
    Color    string
    EnqueueTime time.Time
    StartTime time.Time
}

```

```

func (c *CodePagerDutyJob) Execute() Result {
    // Mock implementation: does nothing and succeeds.
    return Result{ID: c.ID, Ent: c.Entity, Err: nil, Payload: map[string]interface{}{"type": "code", "driver": "pagerduty", "color": c.Color}}
}

```

```

func (c *CodePagerDutyJob) Copy() Job          { job := *c; return &job }
func (c *CodePagerDutyJob) GetEnqueueTime() time.Time { return c.EnqueueTime }
func (c *CodePagerDutyJob) SetEnqueueTime(t time.Time) { c.EnqueueTime = t }
func (c *CodePagerDutyJob) GetStartTime() time.Time { return c.StartTime }
func (c *CodePagerDutyJob) SetStartTime(t time.Time) { c.StartTime = t }

```

```
// CodeSlackJob is a placeholder for a Slack notification job.
```

```

type CodeSlackJob struct {
    ID      uuid.UUID
    Entity   ecs.Entity
    Monitor  string
}

```

```

Message      string
Color        string
EnqueueTime time.Time
StartTime    time.Time
}

func (c *CodeSlackJob) Execute() Result {
    // Mock implementation: does nothing and succeeds.
    return Result{ID: c.ID, Ent: c.Entity, Err: nil, Payload: map[string]interface{}{"type": "code", "driver": "slack", "color": c.Color}}
}

func (c *CodeSlackJob) Copy() Job           { job := *c; return &job }
func (c *CodeSlackJob) GetEnqueueTime() time.Time { return c.EnqueueTime }
func (c *CodeSlackJob) SetEnqueueTime(t time.Time) { c.EnqueueTime = t }
func (c *CodeSlackJob) GetStartTime() time.Time   { return c.StartTime }
func (c *CodeSlackJob) SetStartTime(t time.Time)   { c.StartTime = t }

// CodeEmailJob is a placeholder for an email notification job.
type CodeEmailJob struct {
    ID          uuid.UUID
    Entity      ecs.Entity
    Monitor     string
    Message     string
    Color       string
    EnqueueTime time.Time
    StartTime   time.Time
}
func (c *CodeEmailJob) Execute() Result {
    // Mock implementation: does nothing and succeeds.
    return Result{ID: c.ID, Ent: c.Entity, Err: nil, Payload: map[string]interface{}{"type": "code", "driver": "email", "color": c.Color}}
}

func (c *CodeEmailJob) Copy() Job           { job := *c; return &job }
func (c *CodeEmailJob) GetEnqueueTime() time.Time { return c.EnqueueTime }
func (c *CodeEmailJob) SetEnqueueTime(t time.Time) { c.EnqueueTime = t }
func (c *CodeEmailJob) GetStartTime() time.Time   { return c.StartTime }
func (c *CodeEmailJob) SetStartTime(t time.Time)   { c.StartTime = t }

// CodeWebhookJob is a placeholder for a webhook notification job.
type CodeWebhookJob struct {
    ID          uuid.UUID
    Entity      ecs.Entity
    Monitor     string
    Message     string
    Color       string
    EnqueueTime time.Time
    StartTime   time.Time
}
func (c *CodeWebhookJob) Execute() Result {
    // Mock implementation: does nothing and succeeds.
    return Result{ID: c.ID, Ent: c.Entity, Err: nil, Payload: map[string]interface{}{"type": "code", "driver": "webhook", "color": c.Color}}
}

func (c *CodeWebhookJob) Copy() Job           { job := *c; return &job }
func (c *CodeWebhookJob) GetEnqueueTime() time.Time { return c.EnqueueTime }
func (c *CodeWebhookJob) SetEnqueueTime(t time.Time) { c.EnqueueTime = t }
func (c *CodeWebhookJob) GetStartTime() time.Time   { return c.StartTime }
func (c *CodeWebhookJob) SetStartTime(t time.Time)   { c.StartTime = t }

/* jobs/results.go */

package jobs

import (
    "github.com/google/uuid"
    "github.com/mlange-42/ark/ecs"
)

// Result is a generic structure for returning the outcome of a job.
// It includes the entity it belongs to, any error that occurred, and a flexible payload.
type Result struct {
    ID      uuid.UUID
    Ent    ecs.Entity
    Err    error
}

```

```

        Payload map[string]interface{}}

// Entity returns the entity associated with the result.
func (r *Result) Entity() ecs.Entity {
    return r.Ent
}

// Error returns the error associated with the result, if any.
func (r *Result) Error() error {
    return r.Err
}

/* loader/loader/errors.go */

package loader

/* loader/loader/loader.go */

package loader

import (
    "cpra/internal/loader/schema"
)

type Loader interface {
    Load() error
    GetManifest() schema.Manifest
}

func NewLoader(loaderType string, filename string) Loader {
    switch loaderType {
    case "yaml":
        yamlLoader := NewYamlLoader(filename)
        return yamlLoader

    default:
        yamlLoader := NewYamlLoader(filename)
        return yamlLoader
    }
}

/* loader/loader/loader_test.go */

package loader

import (
    "fmt"
    "testing"
)

func BenchmarkPrimeNumbers(b *testing.B) {
    for i := 0; i < b.N; i++ {
        l := NewYamlLoader("test.yaml")
        l.Load()
        m := l.GetManifest()
        if testing.Verbose() {
            fmt.Printf("loading %d monitors from %s\n", len(m.Monitors), "test.yaml")
        }
    }
}

/* loader/loader/utils.go */

package loader

import (
    "fmt"
    "os"
)

func fatalManifestError(err error) {
    _, err = fmt.Fprintf(os.Stderr, "error: %v\n", err)
    if err != nil {
        return
    }
}

```

```

        }
        os.Exit(1)
    }

/* loader/loader/yaml_loader.go */

package loader

import (
    "cpra/internal/loader/parser"
    "errors"
    "fmt"
    "strings"

    // "cpra/internal/loader/parser"
    "cpra/internal/loader/schema"
    // "errors"
    // "fmt"
    "gopkg.in/yaml.v3"
    // "strings"

    "os"
)

type YamlLoader struct {
    File      string
    Manifest schema.Manifest
}

func NewYamlLoader(fileName string) *YamlLoader {
    return &YamlLoader{
        fileName,
        schema.Manifest{},
    }
}

func (l *YamlLoader) Load() error {
    file, err := os.Open(l.File)
    defer file.Close()

    if err != nil {
        return err
    }
    // decoder := yaml.NewDecoder(file)
    // var manifest schema.Manifest
    // if err := decoder.Decode(&manifest); err != nil {
    //     // This error will now include line numbers and be very clear
    //     // because it comes directly from the yaml.v3 library.
    //     log.Fatal(err)
    //}
    yamlParser := parser.NewYamlParser()
    manifest, err := yamlParser.Parse(file)
    if err != nil {
        var typeErr *yaml.TypeError
        if errors.As(err, &typeErr) {
            for _, msg := range typeErr.Errors {
                if strings.HasPrefix(msg, "line") {
                    return fmt.Errorf("invalid manifest: %s", msg)
                }
            }
        }
        return fmt.Errorf("invalid manifest: %w", err)
    }
    // yamlValidator := validator.NewYamlValidator()
    // err = yamlValidator.ValidateManifest(&manifest)
    // if err != nil {
    //     log.Fatal(err)
    //}
    l.Manifest = manifest
    return nil
}

func (l *YamlLoader) GetManifest() schema.Manifest {
    return l.Manifest
}

/* loader/parser/errors.go */

```

```
package parser

import "fmt"

var (
    ErrInvalidYamlFormat = fmt.Errorf("invalid yaml format")
    ErrUnknownField     = fmt.Errorf("unknown field")
    ErrInvalidPulseType = fmt.Errorf("invalid pulse type")
    ErrRequiredField    = fmt.Errorf("required field")
    ErrInvalidType      = fmt.Errorf("invalid type")
)
type requiredMonitorFieldError struct {
    monitor  string
    parentKey string
    field    string
    line     int
    reason   error
}
func (e *requiredMonitorFieldError) Error() string {
    if e.monitor == "" {
        return fmt.Sprintf("missing required %s field %q (line %d)", e.parentKey, e.field, e.line)
    } else {
        return fmt.Sprintf("missing required %s field %q in monitor %q (line %d)", e.parentKey, e.field, e.monitor, e.line)
    }
}

type monitorFieldTypeError struct {
    FieldName string
    FiledType string
    validType string
}
func (r *monitorFieldTypeError) Error() string {
    return fmt.Sprintf("invalid field type %s for %q: valid types are %s", r.FiledType, r.FieldName, r.validType)
}

type invalidMonitorFieldError struct {
    monitor  string
    parentKey string
    field    string
    line     int
    reason   error
}
func (e *invalidMonitorFieldError) Error() string {
    return fmt.Sprintf("invalid %s field %q in monitor %q (line %d): %s", e.parentKey, e.field, e.monitor, e.line, e.reason)
}

// Unwrap NOT USED REMOVE LATER
func (e *invalidMonitorFieldError) Unwrap() error {
    return e.reason
}

type duplicateMonitorNameError struct {
    name string
    line int
}
func (e *duplicateMonitorNameError) Error() string {
    return fmt.Sprintf("duplicate monitor name %q (line %d), monitor names must be unique and cannot be reused", e.name, e.line)
}

/* loader/parser/parser.go */

package parser

import (
    "cpra/internal/loader/schema"
    "io"
)
```

```

type Parser interface {
    Parse(r io.Reader) (schema.Manifest, error)
}

func NewParser() Parser {
    yamlParser := NewYamlParser()
    return yamlParser
}

/* loader/parser/parser_test.go */

package parser

/* loader/parser/utils.go */

package parser

import (
    "cpra/internal/loader/schema"
    "gopkg.in/yaml.v3"
)

func isValidKey(key string, section string) error {
    f := ManifestFields[section]
    if _, ok := f[key]; !ok {
        return ErrUnknownField
    }
    return nil
}

func checkMissingRequiredKey(section string, node map[string]yaml.Node) (string, error) {
    f := ManifestFields[section]
    for key, field := range f {
        if field.Required {
            if _, ok := node[key]; !ok {
                return key, ErrRequiredField
            }
        }
    }
    return "", nil
}

func decodePulseConfig(config yaml.Node, pulseType string) (schema.PulseConfig, error) {
    switch pulseType {
    case "http":
        var pulseConfig schema.PulseHTTPConfig
        err := config.Decode(&pulseConfig)
        if err != nil {
            return nil, err
        }
        return &pulseConfig, nil
    case "tcp":
        var pulseConfig schema.PulseTCPConfig
        err := config.Decode(&pulseConfig)
        if err != nil {
            return nil, err
        }
        return &pulseConfig, nil
    case "icmp":
        var pulseConfig schema.PulseICMPConfig
        err := config.Decode(&pulseConfig)
        if err != nil {
            return nil, err
        }
        return &pulseConfig, nil
    default:
        return nil, ErrInvalidPulseType
    }
}

/* loader/parser/yaml_parser.go */

package parser

import (
    "cpra/internal/loader/schema"

```

```

"fmt"
"gopkg.in/yaml.v3"
"io"
)

type FieldType struct {
    Required bool
}

type parseState struct {
    line                  int
    monitorName          string
    pulseType             string
    interventionTarget   string
    interventionTargetFields string
    codeColor             string
    fields                string
}

var (
    MonitorFields = map[string]FieldType{
        "name":           {Required: true},
        "enabled":        {Required: false},
        "pulse_check":   {Required: true},
        "intervention":  {Required: false},
        "codes":          {Required: false},
        "notify_groups": {Required: false},
    }
    PulseFields = map[string]FieldType{
        "type":            {Required: true},
        "interval":       {Required: true},
        "timeout":        {Required: true},
        "max_failures":  {Required: false},
        "config":         {Required: true},
    }
    PulseConfigHTTPFields = map[string]FieldType{
        "url":             {Required: true},
        "method":          {Required: false},
        "headers":         {Required: false},
        "auth":            {Required: false},
        "retries":         {Required: false},
    }
    PulseConfigTCPFields = map[string]FieldType{
        "host":            {Required: true},
        "port":            {Required: true},
        "retries":         {Required: false},
    }
    PulseConfigICMPFields = map[string]FieldType{
        "host":            {Required: true},
        "count":           {Required: false},
        "retries":         {Required: false},
        "ignore_privilege": {Required: false},
    }
}

// TODO

//PulseConfigGRPCFields = map[string]FieldType{}
//PulseConfigDockerFields = map[string]FieldType{}

// *****
//PulseConfigTLSFields = map[string]FieldType{}
//PulseConfigUDPFIELDS = map[string]FieldType{}
//PulseConfigDNSFields = map[string]FieldType{}

InterventionFields = map[string]FieldType{
    "action":           {Required: true},
    "retries":          {Required: false},
    "target":           {Required: true},
    "max_failures":    {Required: false},
}
InterventionTargetDockerFields = map[string]FieldType{
    "type":             {Required: false},
    "container":        {Required: true},
    "timeout":          {Required: false},
}
CodeFields = map[string]FieldType{
    "groups":           {Required: false},
    "red":              {Required: false},
}

```

```

        "yellow": {Required: false},
        "green": {Required: false},
        "cyan": {Required: false},
        "gray": {Required: false},
    }

    CodeColorFields = map[string]FieldType{
        "groups": {Required: false},
        "dispatch": {Required: false},
        "notify": {Required: true},
        "config": {Required: true},
    }
}

var ManifestFields = map[string]map[string]FieldType{
    "monitors": MonitorFields,
    "pulse_check": PulseFields,
    "pulse_check_http": PulseConfigHTTPFields,
    "pulse_check_tcp": PulseConfigTCPFields,
    "pulse_check_icmp": PulseConfigICMPFields,
    "intervention": InterventionFields,
    "intervention_docker": InterventionTargetDockerFields,
    "codes": CodeFields,
    "code_color": CodeColorFields,
}
}

type YamlParser struct {}

func NewYamlParser() *YamlParser {
    return &YamlParser{}
}

func (p *YamlParser) Parse(r io.Reader) (schema.Manifest, error) {

    var state parseState
    var manifest schema.Manifest
    decoder := yaml.NewDecoder(r)

    for {
        var node map[string]yaml.Node
        err := decoder.Decode(&node)
        if err == io.EOF {
            break
        }
        if err != nil {
            return schema.Manifest{}, err
        }
        for key, value := range node {
            if key == "monitors" {
                var monitorsNode []yaml.Node
                if err := value.Decode(&monitorsNode); err != nil {
                    return schema.Manifest{}, err
                }
                seen := map[string]struct{}{}

                for _, monitor := range monitorsNode {

                    m, err := p.ParseMonitor(monitor, &state)
                    if err != nil {
                        return schema.Manifest{}, err
                    }
                    name := state.monitorName
                    if _, exists := seen[name]; exists {
                        return schema.Manifest{}, &duplicateMonitorNameError{
                            name: name,
                            line: monitor.Content[0].Line,
                        }
                    }
                    seen[state.monitorName] = struct{}{}

                    manifest.Monitors = append(manifest.Monitors, m)
                }
            }
        }
    }

    return manifest, nil
}

```

```

func (p *YamlParser) ParseMonitor(m yaml.Node, state *parseState) (schema.Monitor, error) {
    var keys map[string]yaml.Node
    monitor := schema.Monitor{}
    if err := m.Decode(&keys); err != nil {
        return schema.Monitor{}, err
    }
    node, ok := keys["name"]
    if !ok {
        node := m.Content[0]
        return schema.Monitor{}, &requiredMonitorFieldError{
            field:      "name",
            parentKey:  "monitor",
            line:       node.Line,
            reason:     ErrRequiredField,
        }
    }
    state.monitorName = node.Value
    state.line = node.Line
    key, err := checkMissingRequiredKey("monitors", keys)

    if err != nil || key != "" {
        return schema.Monitor{}, &requiredMonitorFieldError{
            field:      key,
            parentKey:  "monitor",
            line:       m.Line,
            reason:     ErrRequiredField,
        }
    }

    for k := range keys {
        err := isValidKey(k, "monitors")
        if err != nil {
            line := keys[k].Line
            return schema.Monitor{}, &invalidMonitorFieldError{parentKey: "monitor",
monitor: state.monitorName, field: k, line: line, reason: err}
        }

        monitor.Name = state.monitorName
        if enabled, ok := keys["enabled"]; ok {
            if enabled.Value == "false" {
                monitor.Enabled = false
            } else {
                monitor.Enabled = true
            }
        } else {
            monitor.Enabled = true
        }

        pulseNode := keys["pulse_check"]

        pulse, err := p.ParsePulse(pulseNode, state)

        if err != nil {
            return schema.Monitor{}, err
        }

        monitor.Pulse = pulse

        if interventionNode, ok := keys["intervention"]; ok {
            intervention, err := p.ParseIntervention(interventionNode, state)
            if err != nil {
                return schema.Monitor{}, err
            }
            monitor.Intervention = intervention
        }

        codeNode := keys["codes"]
        codes, err := p.ParseCode(codeNode, state)
        if err != nil {
            return schema.Monitor{}, err
        }
        monitor.Codes = codes

        return monitor, nil
    }

    func (p *YamlParser) ParsePulse(pNode yaml.Node, state *parseState) (schema.Pulse, error) {

```

```

var keys map[string]yaml.Node

if err := pNode.Decode(&keys); err != nil {
    return schema.Pulse{}, err
}

key, err := checkMissingRequiredKey("pulse_check", keys)

if err != nil || key == "" {
    return schema.Pulse{}, &requiredMonitorFieldError{
        field:      key,
        parentKey: "pulse_check",
        line:       pNode.Line,
        reason:     ErrRequiredField,
    }
}

for k := range keys {
    err := isValidKey(k, "pulse_check")
    if err != nil {
        line := keys[k].Line

        return schema.Pulse{}, &invalidMonitorFieldError{parentKey: "pulse_check",
monitor: state.monitorName, field: k, line: line, reason: err}
    }
}

state.pulseType = keys["type"].Value

switch state.pulseType {
case "http":
    state.fields = "pulse_check_http"
case "tcp":
    state.fields = "pulse_check_tcp"
case "icmp":
    state.fields = "pulse_check_icmp"

default:
    return schema.Pulse{}, &invalidMonitorFieldError{parentKey: "pulse_check", monitor:
state.monitorName, field: "type", line: keys["type"].Line, reason: ErrUnknownField}
}

state.line = pNode.Content[0].Line
pConfig := keys["config"]
config, err := p.ParsePulseConfig(pConfig, state)
if err != nil {
    return schema.Pulse{}, err
}

var pulse schema.Pulse
err = pNode.Decode(&pulse)
if err != nil {
    return schema.Pulse{}, err
}
pulse.Config = config
return pulse, nil
}

func (p *YamlParser) ParsePulseConfig(pNode yaml.Node, state *parseState) (schema.PulseConfig, error)
{
    var keys map[string]yaml.Node
    if err := pNode.Decode(&keys); err != nil {
        return nil, err
    }
    state.line = pNode.Content[0].Line

    key, err := checkMissingRequiredKey(state.fields, keys)

    if err != nil || key == "" {
        return nil, &requiredMonitorFieldError{
            field:      key,
            parentKey: fmt.Sprintf("%v pulse_check", state.pulseType),
            line:       state.line,
            reason:     ErrRequiredField,
        }
    }

    for k := range keys {
        err := isValidKey(k, state.fields)
        if err != nil {

```

```

        line := keys[k].Line
        return nil, &invalidMonitorFieldError{parentKey: "pulse_check", monitor:
state.monitorName, field: k, line: line, reason: fmt.Errorf("invalid pulse config for type %q %w",
keys["type"].Value, err)}
    }
}

pulseType, err := decodePulseConfig(pNode, state.pulseType)

if err != nil {
    return nil, err
}

return pulseType, nil
}

func (p *YamlParser) ParseIntervention(i yaml.Node, state *parseState) (schema.Intervention, error) {
var keys map[string]yaml.Node
if err := i.Decode(&keys); err != nil {
    return schema.Intervention{}, err
}

key, err := checkMissingRequiredKey("intervention", keys)

if err != nil || key == "" {
    return schema.Intervention{}, &requiredMonitorFieldError{
        field:      key,
        parentKey: "intervention",
        line:       i.Line,
        reason:     ErrRequiredField,
    }
}

for k := range keys {
    err := isValidKey(k, "intervention")
    if err != nil {
        line := keys[k].Line

        return schema.Intervention{}, &invalidMonitorFieldError{parentKey:
"intervention", monitor: state.monitorName, field: k, line: line, reason: err}
    }
}

iTargt := keys["target"]

var target map[string]yaml.Node
if err := iTargt.Decode(&target); err != nil {
    return schema.Intervention{}, err
}

//targetType, ok := target["type"]
//if !ok {
//    return schema.Intervention{}, &requiredMonitorFieldError{
//        field:      "type",
//        parentKey: "intervention.target",
//        line:       targetType.Line,
//        reason:     ErrRequiredField,
//    }
//}
action := keys["action"]
switch action.Value {
case "docker":
    state.interventionTargetFields = "intervention_docker"
    err := p.ParseInterventionTarget(iTargt, state)
    if err != nil {
        return schema.Intervention{}, err
    }
default:
    return schema.Intervention{}, &invalidMonitorFieldError{parentKey: "intervention",
monitor: state.monitorName, field: "target", line: keys["target"].Content[0].Line, reason:
ErrUnknownField}
}

var intervention schema.Intervention
err = i.Decode(&intervention)
if err != nil {
    return schema.Intervention{}, err
}

//pulse.Config = config
return intervention, nil
}

```

```

func (p *YamlParser) ParseInterventionTarget(i yaml.Node, state *parseState) error {
    var keys map[string]yaml.Node

    if err := i.Decode(&keys); err != nil {
        return err
    }

    key, err := checkMissingRequiredKey(state.interventionTargetFields, keys)

    if err != nil || key != "" {
        return &requiredMonitorFieldError{
            field:      key,
            parentKey: fmt.Sprintf("%v intervention", keys["type"].Value),
            line:       i.Line,
            reason:     ErrRequiredField,
        }
    }

    for k := range keys {
        err := isValidKey(k, state.interventionTargetFields)
        if err != nil {
            line := keys[k].Line
            return &invalidMonitorFieldError{parentKey: fmt.Sprintf("%v intervention",
keys["type"].Value), monitor: state.monitorName, field: k, line: line, reason: err}
        }
    }
    return nil
}

func (p *YamlParser) ParseCode(c yaml.Node, state *parseState) (schema.Codes, error) {
    var keys map[string]yaml.Node
    if err := c.Decode(&keys); err != nil {
        return nil, err
    }

    var codes schema.Codes
    if err := c.Decode(&codes); err != nil {
        return nil, err
    }

    key, err := checkMissingRequiredKey("codes", keys)
    if err != nil || key != "" {
        return nil, &requiredMonitorFieldError{
            field:      key,
            parentKey: "codes",
            line:       c.Line,
            reason:     ErrRequiredField,
        }
    }

    //codes := make(schema.Codes)

    for k, _ := range keys {
        err := isValidKey(k, "codes")
        if err != nil {
            line := keys[k].Line
            return nil, &invalidMonitorFieldError{parentKey: "codes", monitor:
state.monitorName, field: k, line: line, reason: err}
        }
    }

    for colorKey, colorNode := range keys {
        // ... (existing key validation)

        // ADD THIS LOGIC
        state.codeColor = colorKey // Set the state for the color being parsed
        err := p.ParseCodeColor(colorNode, state)
        if err != nil {
            return nil, err
        }
    }
    return codes, nil
}

func (p *YamlParser) ParseCodeColor(c yaml.Node, state *parseState) error {
    var keys map[string]yaml.Node
    if err := c.Decode(&keys); err != nil {
        return err
    }

    key, err := checkMissingRequiredKey("code_color", keys)

```

```

if err != nil || key == "" {
    return &requiredMonitorFieldError{
        field:      key,
        parentKey: state.codeColor,
        line:       c.Line,
        reason:     ErrRequiredField,
    }
}

for k := range keys {
    err := isValidKey(k, "code_color")
    if err != nil {
        line := keys[k].Line
        return &invalidMonitorFieldError{parentKey: state.codeColor, monitor:
state.monitorName, field: k, line: line, reason: err}
    }
}
return nil
}

/* loader/schema/manifest.go */

package schema

import (
    "encoding/json"
    "fmt"
    "strings"
    "time"

    "gopkg.in/yaml.v3"
)

//// UTILITY TYPES

type DurationSeconds int

func (d *DurationSeconds) UnmarshalYAML(unmarshal func(interface{}) error) error {
    var raw string
    if err := unmarshal(&raw); err != nil {
        return fmt.Errorf("invalid duration %q", raw)
    }
    p, err := time.ParseDuration(raw)
    if err != nil {
        return fmt.Errorf("invalid duration %q: %w", raw, err)
    }
    *d = DurationSeconds(int(p.Seconds()))
    return nil
}

type StringList []string

func (s *StringList) UnmarshalYAML(unmarshal func(interface{}) error) error {
    var single string
    if err := unmarshal(&single); err == nil {
        *s = []string{single}
        return nil
    }
    var multi []string
    if err := unmarshal(&multi); err == nil {
        *s = multi
        return nil
    }
    return fmt.Errorf("value must be a string or list of strings")
}

//// PULSE TYPES

type PulseConfig interface {
    isPulseConfigs()
    Copy() PulseConfig
}

type PulseHTTPConfig struct {
    Url      string      `yaml:"url" json:"url"`
    Method   string      `yaml:"method" json:"method"`
    Headers  StringList  `yaml:"headers" json:"headers"`
    Retries  int         `yaml:"retries" json:"retries"`
}

```

```

}

func (c *PulseHTTPConfig) Copy() PulseConfig {
    // This was already correct, but for consistency, we'll return a pointer
    // to a new struct.

    newConfig := new(PulseHTTPConfig)
    *newConfig = *c
    return newConfig
}

func (*PulseHTTPConfig) isPulseConfigs() {}

type PulseTCPConfig struct {
    Host      string `yaml:"host"`
    Port      int    `yaml:"port"`
    Retries   int    `yaml:"retries"`
}

func (c *PulseTCPConfig) Copy() PulseConfig {
    newConfig := new(PulseTCPConfig)
    *newConfig = *c
    return newConfig
}

func (*PulseTCPConfig) isPulseConfigs() {}

type PulseICMPConfig struct {
    Host      string `yaml:"host"`
    Privilege bool   `yaml:"ignore_privilege"`
    Count     int    `yaml:"count"`
    Retries   int    `yaml:"retries"`
}

func (c *PulseICMPConfig) Copy() PulseConfig {
    newConfig := new(PulseICMPConfig)
    *newConfig = *c
    return newConfig
}

func (*PulseICMPConfig) isPulseConfigs() {}

type Pulse struct {
    Type          string      `yaml:"type" json:"type"`
    Interval      time.Duration `yaml:"interval" json:"interval"`
    Timeout       time.Duration `yaml:"timeout" json:"timeout"`
    MaxFailures  int         `yaml:"max_failures" json:"max_failures"`
    Groups        StringList   `yaml:"groups" json:"groups"`
    Config        PulseConfig  `json:"config"`
}

type rawPulse struct {
    Type          string      `yaml:"type"`
    Interval      time.Duration `yaml:"interval"`
    Timeout       time.Duration `yaml:"timeout"`
    Retries       int         `yaml:"retries"`
    MaxFailures  int         `yaml:"max_failures"`
    Groups        StringList   `yaml:"groups"`
}

func (p *Pulse) UnmarshalYAML(value *yaml.Node) error {
    var temp struct {
        Config  yaml.Node `yaml:"config"`
        rawPulse `yaml:",inline"`
    }
    if err := value.Decode(&temp); err != nil {
        return err
    }
    *p = Pulse{
        Type:      temp.Type,
        Interval:  temp.Interval,
        Timeout:   temp.Timeout,
        MaxFailures: temp.MaxFailures,
        Groups:    temp.Groups,
    }
    switch temp.Type {
    case "http":
        var c = &PulseHTTPConfig{} // FIX: Allocate on the heap
        if err := temp.Config.Decode(c); err != nil {
            return err
        }
    }
}

```

```

        }
        p.Config = c
    case "tcp":
        var c = &PulseTCPConfig{} // FIX: Allocate on the heap
        if err := temp.Config.Decode(c); err != nil {
            return err
        }
        p.Config = c
    case "icmp":
        var c = &PulseICMPConfig{} // FIX: Allocate on the heap
        if err := temp.Config.Decode(c); err != nil {
            return err
        }
        p.Config = c
    default:
        return fmt.Errorf("unknown pulse type: %q", temp.Type)
    }
    return nil
}

// UnmarshalJSON handles JSON unmarshaling for Pulse (needed for JSON parser)
func (p *Pulse) UnmarshalJSON(data []byte) error {
    var temp struct {
        Type      string      `json:"type"`
        Interval  string      `json:"interval"`   // Parse as string first
        Timeout   string      `json:"timeout"`   // Parse as string first
        MaxFailures int       `json:"max_failures"`
        Config    json.RawMessage `json:"config"`
    }
    if err := json.Unmarshal(data, &temp); err != nil {
        return err
    }

    // Parse duration strings
    interval, err := time.ParseDuration(temp.Interval)
    if err != nil {
        return fmt.Errorf("invalid interval duration %q: %w", temp.Interval, err)
    }

    timeout, err := time.ParseDuration(temp.Timeout)
    if err != nil {
        return fmt.Errorf("invalid timeout duration %q: %w", temp.Timeout, err)
    }

    *p = Pulse{
        Type:      temp.Type,
        Interval:  interval,
        Timeout:   timeout,
        MaxFailures: temp.MaxFailures,
    }

    switch temp.Type {
    case "http":
        var c = &PulseHTTPConfig{}
        if err := json.Unmarshal(temp.Config, c); err != nil {
            return err
        }
        p.Config = c
    case "tcp":
        var c = &PulseTCPConfig{}
        if err := json.Unmarshal(temp.Config, c); err != nil {
            return err
        }
        p.Config = c
    case "icmp":
        var c = &PulseICMPConfig{}
        if err := json.Unmarshal(temp.Config, c); err != nil {
            return err
        }
        p.Config = c
    default:
        return fmt.Errorf("unknown pulse type: %q", temp.Type)
    }
    return nil
}

//// INTERVENTION TYPES

type Intervention struct {

```

```

        Action      string      `yaml:"action"`
        Retries    int         `yaml:"retries"`
        Target     InterventionTarget `yaml:"target"`
        MaxFailures int       `yaml:"max_failures"`
    }
}

type rawIntervention struct {
    Action string `yaml:"action"`
    Retries int   `yaml:"retries"`
}

func (i *Intervention) UnmarshalYAML(value *yaml.Node) error {
    var temp struct {
        Target     yaml.Node `yaml:"target"`
        rawIntervention `yaml:",inline"`
    }
    if err := value.Decode(&temp); err != nil {
        return err
    }
    *i = Intervention{
        Action:  temp.Action,
        Retries: temp.Retries,
    }
    switch temp.Action {
    case "docker":
        var t = &InterventionTargetDocker{} // FIX: Allocate on the heap
        if err := temp.Target.Decode(t); err != nil {
            return err
        }
        i.Target = t
    default:
        return fmt.Errorf("unknown intervention type: %q", temp.Action)
    }
    return nil
}

// UnmarshalJSON handles JSON unmarshaling for Intervention (needed for JSON parser)
func (i *Intervention) UnmarshalJSON(data []byte) error {
    var temp struct {
        Action string      `json:"action"`
        Retries int         `json:"retries"`
        Target json.RawMessage `json:"target"`
    }

    if err := json.Unmarshal(data, &temp); err != nil {
        return err
    }

    *i = Intervention{
        Action:  temp.Action,
        Retries: temp.Retries,
    }

    switch temp.Action {
    case "docker":
        var t = &InterventionTargetDocker{}
        if err := json.Unmarshal(temp.Target, t); err != nil {
            return err
        }
        i.Target = t
    default:
        return fmt.Errorf("unknown intervention type: %q", temp.Action)
    }
    return nil
}

type InterventionTarget interface {
    GetTargetType() string
    Copy() InterventionTarget
}

type InterventionTargetDocker struct {
    Type      string      `yaml:"type" json:"type"`
    Container string      `yaml:"container" json:"container"`
    Timeout   time.Duration `yaml:"timeout" json:"timeout"`
}

func (i *InterventionTargetDocker) Copy() InterventionTarget {
    return &InterventionTargetDocker{
        Type:      strings.Clone(i.Type),
        Container: strings.Clone(i.Container),
    }
}

```

```

}

func (i *InterventionTargetDocker) GetTargetType() string {
    return i.Type
}

type CodeNotification interface {
    IsCodeNotification()
    Copy() CodeNotification
}

type CodeNotificationLog struct {
    File string `yaml:"file" json:"file"`
}

func (c *CodeNotificationLog) Copy() CodeNotification {
    return &CodeNotificationLog{
        File: strings.Clone(c.File),
    }
}

func (c *CodeNotificationLog) IsCodeNotification() {
}

type CodeNotificationPagerDuty struct {
    URL string `yaml:"url" json:"url"`
}

func (c *CodeNotificationPagerDuty) Copy() CodeNotification {
    return &CodeNotificationPagerDuty{
        URL: strings.Clone(c.URL),
    }
}

func (c *CodeNotificationPagerDuty) IsCodeNotification() {
}

type CodeNotificationSlack struct {
    WebHook string `yaml:"hook" json:"hook"`
}

func (c *CodeNotificationSlack) Copy() CodeNotification {
    return &CodeNotificationSlack{
        WebHook: strings.Clone(c.WebHook),
    }
}

func (c *CodeNotificationSlack) IsCodeNotification() {
}

type CodeConfig struct {
    Dispatch bool           `yaml:"dispatch"`
    Notify   string          `yaml:"notify"`
    Config   CodeNotification `yaml:"config"` // or more specific struct if desired
}

type Codes map[string]CodeConfig

type rawCodes struct {
    Dispatch bool `yaml:"dispatch"`
    Notify   string `yaml:"notify"`
}

func (c *Codes) UnmarshalYAML(value *yaml.Node) error {
    var codes map[string]yaml.Node
    if err := value.Decode(&codes); err != nil {
        return err
    }
    colors := make(map[string]CodeConfig)
    for color, config := range codes {
        var temp struct {
            Config   yaml.Node `yaml:"config"`
            rawCodes `yaml:",inline"`
        }
        if err := config.Decode(&temp); err != nil {
            return err
        }
        switch temp.Notify {
        case "log":

```

```

        var t = &CodeNotificationLog{} // FIX: Allocate on the heap
        if err := temp.Config.Decode(t); err != nil {
            return err
        }
        colors[color] = CodeConfig{
            Dispatch: temp.Dispatch,
            Notify:   temp.Notify,
            Config:   t,
        }
    case "slack":
        var t = &CodeNotificationSlack{} // FIX: Allocate on the heap
        if err := temp.Config.Decode(t); err != nil {
            return err
        }
        colors[color] = CodeConfig{
            Dispatch: temp.Dispatch,
            Notify:   temp.Notify,
            Config:   t,
        }
    case "pagerduty":
        var t = &CodeNotificationPagerDuty{} // FIX: Allocate on the heap
        if err := temp.Config.Decode(t); err != nil {
            return err
        }
        colors[color] = CodeConfig{
            Dispatch: temp.Dispatch,
            Notify:   temp.Notify,
            Config:   t,
        }
    default:
        return fmt.Errorf("unknown notification type: %q", temp.Notify)
    }
}
*c = colors
return nil
}

// UnmarshalJSON handles JSON unmarshaling for Codes (needed for JSON parser)
func (c *Codes) UnmarshalJSON(data []byte) error {
    var codes map[string]struct {
        Dispatch bool           `json:"dispatch"`
        Notify   string          `json:"notify"`
        Config   json.RawMessage `json:"config"`
    }

    if err := json.Unmarshal(data, &codes); err != nil {
        return err
    }

    colors := make(map[string]CodeConfig)
    for color, config := range codes {
        switch config.Notify {
        case "log":
            var t = &CodeNotificationLog{}
            if err := json.Unmarshal(config.Config, t); err != nil {
                return err
            }
            colors[color] = CodeConfig{
                Dispatch: config.Dispatch,
                Notify:   config.Notify,
                Config:   t,
            }
        case "slack":
            var t = &CodeNotificationSlack{}
            if err := json.Unmarshal(config.Config, t); err != nil {
                return err
            }
            colors[color] = CodeConfig{
                Dispatch: config.Dispatch,
                Notify:   config.Notify,
                Config:   t,
            }
        case "pagerduty":
            var t = &CodeNotificationPagerDuty{}
            if err := json.Unmarshal(config.Config, t); err != nil {
                return err
            }
            colors[color] = CodeConfig{
                Dispatch: config.Dispatch,
                Notify:   config.Notify,
                Config:   t,
            }
        }
    }
}

```

```

        Config:  t,
    }
    default:
        return fmt.Errorf("unknown notification type: %q", config.Notify)
    }
}
*c = colors
return nil
}

type Monitor struct {
    Name          string      `yaml:"name" json:"name"`
    Enabled       bool        `yaml:"enabled" json:"enabled"`
    Pulse         Pulse       `yaml:"pulse_check" json:"pulse_check"`
    Intervention Intervention `yaml:"intervention,omitempty" json:"intervention,omitempty"`
    Codes         Codes       `yaml:"codes" json:"codes"`
}

// UnmarshalYAML sets default values for the Monitor struct, specifically for the Enabled field.
func (m *Monitor) UnmarshalYAML(value *yaml.Node) error {
    // Create a temporary struct with a pointer to a bool for 'Enabled'
    type TmpMonitor struct {
        Name          string      `yaml:"name"`
        Enabled       *bool       `yaml:"enabled"`
        Pulse         Pulse       `yaml:"pulse_check"`
        Intervention Intervention `yaml:"intervention,omitempty"`
        Codes         Codes       `yaml:"codes"`
    }

    var tmp TmpMonitor
    if err := value.Decode(&tmp); err != nil {
        return err
    }

    // Assign fields to the actual monitor struct
    m.Name = tmp.Name
    m.Pulse = tmp.Pulse
    m.Intervention = tmp.Intervention
    m.Codes = tmp.Codes

    // Set 'Enabled' to true if it's not specified in the YAML
    if tmp.Enabled == nil {
        m.Enabled = true
    } else {
        m.Enabled = *tmp.Enabled
    }

    return nil
}

type Manifest struct {
    Monitors []Monitor `yaml:"monitors" json:"monitors"`
}

/*
 * loader/streaming/streaming_entity_creator.go */
package streaming

import (
    "context"
    "cpra/internal/controller/entities"
    "fmt"
    "sync"
    "time"

    "github.com/mlange-42/ark/ecs"
)

// StreamingEntityCreator handles batch entity creation for Ark ECS.
// It now uses the consolidated EntityManager to create entities.
type StreamingEntityCreator struct {
    world          *ecs.World
    entityManager *entities.EntityManager

    // Statistics
    entitiesCreated int64
    batchesProcessed int64
    startTime       time.Time
    pulseRate       float64
}

```

```

    mu          sync.RWMutex
}

// EntityCreationConfig holds entity creation configuration.
type EntityCreationConfig struct {
    BatchSize    int
    PreAllocate  int
    ProgressChan chan<- EntityProgress
}

// EntityProgress represents entity creation progress.
type EntityProgress struct {
    EntitiesCreated int64
    BatchesProcessed int64
    Rate            float64
    MemoryUsage     int64
}

// NewStreamingEntityCreator creates a new, simplified entity creator.
func NewStreamingEntityCreator(world *ecs.World, config EntityCreationConfig) *StreamingEntityCreator {
    creator := &StreamingEntityCreator{
        world:           world,
        entityManager: entities.NewEntityManager(world),
        startTime:      time.Now(),
    }

    if config.PreAllocate > 0 {
        creator.preAllocateEntities(config.PreAllocate)
    }

    return creator
}

// preAllocateEntities pre-allocates entity storage to reduce allocations during creation.
func (c *StreamingEntityCreator) preAllocateEntities(count int) {
    tempEntities := make([]ecs.Entity, count)
    for i := 0; i < count; i++ {
        tempEntities[i] = c.world.NewEntity()
    }
    for _, entity := range tempEntities {
        c.world.RemoveEntity(entity)
    }
}

// ProcessBatches processes monitor batches and creates entities.
func (c *StreamingEntityCreator) ProcessBatches(ctx context.Context, batchChan <-chan MonitorBatch,
progressChan chan<- EntityProgress) error {
    progressTicker := time.NewTicker(2 * time.Second)
    defer progressTicker.Stop()

    for {
        select {
        case <-ctx.Done():
            return ctx.Err()

        case <-progressTicker.C:
            c.reportProgress(progressChan)

        case batch, ok := <-batchChan:
            if !ok {
                c.reportProgress(progressChan)
                return nil
            }

            if err := c.processBatch(batch); err != nil {
                return fmt.Errorf("failed to process batch %d: %w", batch.BatchID,
err)
            }
        }
    }
}

// processBatch creates entities for a single batch of monitors.
func (c *StreamingEntityCreator) processBatch(batch MonitorBatch) error {
    var pulseSum float64
    for _, monitor := range batch.Monitors {
        if err := c.entityManager.CreateEntityFromMonitor(&monitor, c.world); err != nil {
            return fmt.Errorf("failed to create entity for monitor '%s': %w",
monitor.Name, err)
        }
    }
}

```

```

        }
        if monitor.Enabled && monitor.Pulse.Interval > 0 {
            sec := monitor.Pulse.Interval.Seconds()
            if sec > 0 {
                pulseSum += 1.0 / sec
            }
        }
    }

    c.mu.Lock()
    c.entitiesCreated += int64(len(batch.Monitors))
    c.batchesProcessed++
    c.pulseRate += pulseSum
    c.mu.Unlock()

    return nil
}

// reportProgress sends a progress update.
func (c *StreamingEntityCreator) reportProgress(progressChan chan<- EntityProgress) {
    if progressChan == nil {
        return
    }

    c.mu.RLock()
    entitiesCreated := c.entitiesCreated
    batchesProcessed := c.batchesProcessed
    c.mu.RUnlock()

    elapsed := time.Since(c.startTime)
    rate := 0.0
    if elapsed.Seconds() > 0 {
        rate = float64(entitiesCreated) / elapsed.Seconds()
    }

    select {
    case progressChan <- EntityProgress{
        EntitiesCreated: entitiesCreated,
        BatchesProcessed: batchesProcessed,
        Rate: rate,
    }:
    default:
    }
}

// PulseRate returns the aggregated expected pulse arrival rate (jobs/sec).
func (c *StreamingEntityCreator) PulseRate() float64 {
    c.mu.RLock()
    defer c.mu.RUnlock()
    return c.pulseRate
}

// GetStats returns current creation statistics.
func (c *StreamingEntityCreator) GetStats() (entitiesCreated int64, batchesProcessed int64, rate float64) {
    c.mu.RLock()
    defer c.mu.RUnlock()

    elapsed := time.Since(c.startTime)
    if elapsed.Seconds() == 0 {
        return c.entitiesCreated, c.batchesProcessed, 0
    }
    return c.entitiesCreated, c.batchesProcessed, float64(c.entitiesCreated) / elapsed.Seconds()
}

/*
 * loader/streaming/streaming_json_parser.go */
package streaming

import (
    "context"
    "encoding/json"
    "fmt"
    "io"
    "os"
)

```

```

// StreamingJsonParser handles true streaming parsing of a JSON file.
// It reads the file object by object, creating batches without loading the entire file into memory.
type StreamingJsonParser struct {
    filename string
    config   ParseConfig
}

// NewStreamingJsonParser creates a new streaming JSON parser.
func NewStreamingJsonParser(filename string, config ParseConfig) (*StreamingJsonParser, error) {
    return &StreamingJsonParser{
        filename: filename,
        config:   config,
    }, nil
}

// ParseBatches streams the JSON file and sends batches of monitors over a channel.
func (p *StreamingJsonParser) ParseBatches(ctx context.Context, progressChan chan<- Progress) (<-chan
MonitorBatch, <-chan error) {
    batchChan := make(chan MonitorBatch, 100) // Buffer for a few batches
    errorChan := make(chan error, 1)

    go func() {
        defer close(batchChan)
        defer close(errorChan)

        if err := p.parseFile(ctx, batchChan, progressChan); err != nil {
            errorChan <- err
        }
    }()
    return batchChan, errorChan
}

// parseFile performs the actual streaming JSON parsing.
func (p *StreamingJsonParser) parseFile(ctx context.Context, batchChan chan<- MonitorBatch,
progressChan chan<- Progress) error {
    file, err := os.Open(p.filename)
    if err != nil {
        return fmt.Errorf("failed to open file: %w", err)
    }
    defer file.Close()

    decoder := json.NewDecoder(file)

    // Read the opening bracket of the object.
    t, err := decoder.Token()
    if err != nil {
        return fmt.Errorf("failed to read opening token: %w", err)
    }
    if t != json.Delim('{') {
        return fmt.Errorf("expected { at start of json file, got %v", t)
    }

    // Find the "monitors" key
    for decoder.More() {
        t, err := decoder.Token()
        if err != nil {
            return fmt.Errorf("failed to read token: %w", err)
        }
        if s, ok := t.(string); ok && s == "monitors" {
            break
        }
    }

    // Read the opening bracket of the array.
    t, err = decoder.Token()
    if err != nil {
        return fmt.Errorf("failed to read opening token: %w", err)
    }
    if t != json.Delim('[') {
        return fmt.Errorf("expected [ after 'monitors' key, got %v", t)
    }

    batchID := 0
    // Loop while there are more objects in the array.
    for decoder.More() {
        select {
        case <-ctx.Done():
            return ctx.Err()
        default:

```

```

batch := make([]schema.Monitor, 0, p.config.BatchSize)
// Fill a batch
for i := 0; i < p.config.BatchSize && decoder.More(); i++ {
    var monitor schema.Monitor
    if err := decoder.Decode(&monitor); err != nil {
        return fmt.Errorf("failed to decode monitor object: %w", err)
    }
    batch = append(batch, monitor)
}

// Send the batch to the creator
batchChan <- MonitorBatch{
    Monitors: batch,
    BatchID:  batchID,
}
batchID++
}

// Read the closing bracket of the array.
t, err = decoder.Token()
if err != nil && err != io.EOF {
    return fmt.Errorf("failed to read closing token: %w", err)
}
if t != json.Delim(']') {
    return fmt.Errorf("expected ] at end of json file, got %v", t)
}

return nil
}

/* loader/streaming/streaming_loader.go */
package streaming

import (
    "context"
    "cpra/internal/loader/schema"
    "fmt"
    "runtime"
    "strings"
    "time"

    "github.com/mlange-42/ark/ecs"
)

// MonitorBatch represents a batch of monitors read from a file.
type MonitorBatch struct {
    Monitors []schema.Monitor
    BatchID  int
    Offset   int64
}

// ParseConfig holds configuration for the streaming parsers.
type ParseConfig struct {
    BatchSize    int
    BufferSize   int
    MaxMemory   int64
    ProgressChan chan<- Progress
}

// Progress represents parsing progress.
type Progress struct {
    EntitiesProcessed int64
    TotalBytes       int64
    ProcessedBytes   int64
    Percentage       float64
    Rate             float64 // entities per second
    EstimatedRemaining time.Duration
}

// StreamingLoader orchestrates the streaming loading process
type StreamingLoader struct {
    filename string
    world    *ecs.World
    config   StreamingConfig

    parseProgress chan Progress
    entityProgress chan EntityProgress
}

```

```

        totalStartTime time.Time
        loadingStats LoadingStats
    }

// StreamingConfig holds all streaming configuration
type StreamingConfig struct {
    ParseBatchSize int
    ParseBufferSize int
    MaxParseMemory int64
    EntityBatchSize int
    PreAllocateCount int
    MaxWorkers int
    ProgressInterval time.Duration
    GCInterval time.Duration
    MemoryLimit int64
}

// LoadingStats holds comprehensive loading statistics
type LoadingStats struct {
    TotalEntities int64
    LoadingTime time.Duration
    ParseRate float64
    CreationRate float64
    MemoryUsage int64
    GCCount int
    PulseRate float64
}

// DefaultStreamingConfig returns optimized default configuration for large files
func DefaultStreamingConfig() StreamingConfig {
    return StreamingConfig{
        ParseBatchSize: 10000,
        ParseBufferSize: 4 * 1024 * 1024,
        MaxParseMemory: 1 * 1024 * 1024 * 1024,
        EntityBatchSize: 10000,
        PreAllocateCount: 500000,
        MaxWorkers: runtime.NumCPU() * 2,
        ProgressInterval: 1 * time.Second,
        GCInterval: 5 * time.Second,
        MemoryLimit: 2 * 1024 * 1024 * 1024,
    }
}

// NewStreamingLoader creates a new streaming loader
func NewStreamingLoader(filename string, world *ecs.World, config StreamingConfig) *StreamingLoader {
    return &StreamingLoader{
        filename: filename,
        world: world,
        config: config,
        parseProgress: make(chan Progress, 10),
        entityProgress: make(chan EntityProgress, 10),
        totalStartTime: time.Now(),
    }
}

// Load performs the complete streaming load operation
func (sl *StreamingLoader) Load(ctx context.Context) (*LoadingStats, error) {
    fmt.Printf("Starting streaming load of %s...\n", sl.filename)

    var batchChan <-chan MonitorBatch
    var errorChan <-chan error

    parseConfig := ParseConfig{
        BatchSize: sl.config.ParseBatchSize,
        BufferSize: sl.config.ParseBufferSize,
        MaxMemory: sl.config.MaxParseMemory,
        ProgressChan: sl.parseProgress,
    }

    if strings.HasSuffix(strings.ToLower(sl.filename), ".json") {
        jsonParser, err := NewStreamingJsonParser(sl.filename, parseConfig)
        if err != nil {
            return nil, fmt.Errorf("failed to create JSON parser: %w", err)
        }
        batchChan, errorChan = jsonParser.ParseBatches(ctx, sl.parseProgress)
    } else {
        yamlParser, err := NewStreamingYamlParser(sl.filename, parseConfig)
        if err != nil {
            return nil, fmt.Errorf("failed to create YAML parser: %w", err)
        }
        batchChan, errorChan = yamlParser.ParseBatches(ctx, sl.parseProgress)
    }

    go func() {
        for batch := range batchChan {
            sl.entityProgress(&batch)
        }
    }()
    go func() {
        for err := range errorChan {
            sl.error(err)
        }
    }()
}

```

```

        }
        batchChan, errorChan = yamlParser.ParseBatches(ctx, sl.parseProgress)
    }

    entityCreator := NewStreamingEntityCreator(sl.world, EntityCreationConfig{
        BatchSize:   sl.config.EntityBatchSize,
        PreAllocate: sl.config.PreAllocateCount,
        ProgressChan: sl.entityProgress,
    })

    err := entityCreator.ProcessBatches(ctx, batchChan, sl.entityProgress)
    if err != nil {
        return nil, fmt.Errorf("failed to create entities: %w", err)
    }

    select {
    case parseErr := <-errorChan:
        if parseErr != nil {
            return nil, fmt.Errorf("parsing error: %w", parseErr)
        }
    default:
    }

    sl.finalizeStats(entityCreator)

    fmt.Printf("Streaming load completed: %d entities in %v (%.0f entities/sec)\n",
        sl.loadingStats.TotalEntities,
        sl.loadingStats.LoadingTime,
        sl.loadingStats.CreationRate)

    return &sl.loadingStats, nil
}

func (sl *StreamingLoader) finalizeStats(creator *StreamingEntityCreator) {
    sl.loadingStats.LoadingTime = time.Since(sl.totalStartTime)
    entitiesCreated, _, creationRate := creator.GetStats()
    sl.loadingStats.TotalEntities = entitiesCreated
    sl.loadingStats.CreationRate = creationRate
    sl.loadingStats.PulseRate = creator.PulseRate()
    var m runtime.MemStats
    runtime.ReadMemStats(&m)
    sl.loadingStats.MemoryUsage = int64(m.Alloc)
}

/* loader/streaming/streaming_yaml_parser.go */
package streaming

import (
    "context"
    "fmt"
    "io"
    "os"

    "cpra/internal/loader/schema"
    "gopkg.in/yaml.v3"
)

// StreamingYamlParser handles true streaming parsing of a YAML file.
// It reads the file document by document, creating batches without loading the entire file into
// memory.
type StreamingYamlParser struct {
    filename string
    config   ParseConfig
}

// NewStreamingYamlParser creates a new streaming YAML parser.
func NewStreamingYamlParser(filename string, config ParseConfig) (*StreamingYamlParser, error) {
    return &StreamingYamlParser{
        filename: filename,
        config:   config,
    }, nil
}

// ParseBatches streams the YAML file and sends batches of monitors over a channel.
func (p *StreamingYamlParser) ParseBatches(ctx context.Context, progressChan chan<- Progress) (<-chan
MonitorBatch, <-chan error) {
    batchChan := make(chan MonitorBatch, 100)
    errorChan := make(chan error, 1)

```

```

go func() {
    defer close(batchChan)
    defer close(errorChan)

    if err := p.parseFile(ctx, batchChan, progressChan); err != nil {
        errorChan <- err
    }
}()

return batchChan, errorChan
}

// parseFile performs the actual streaming YAML parsing.
// It decodes the main `monitors` list and then streams each monitor entry.
func (p *StreamingYamlParser) parseFile(ctx context.Context, batchChan chan<- MonitorBatch,
progressChan chan<- Progress) error {
    file, err := os.Open(p.filename)
    if err != nil {
        return fmt.Errorf("failed to open file: %w", err)
    }
    defer file.Close()

    decoder := yaml.NewDecoder(file)

    // The YAML file is expected to have a root structure like:
    // monitors:
    //   - name: ...
    //   - name: ...
    // We need to find the 'monitors' sequence node.

    // 1. Decode the top-level structure.
    var topLevel struct {
        Monitors yaml.Node `yaml:"monitors"`
    }
    if err := decoder.Decode(&topLevel); err != nil {
        if err == io.EOF {
            return nil // Empty file is not an error
        }
        return fmt.Errorf("failed to decode top-level 'monitors' field: %w", err)
    }

    // 2. Check if 'monitors' is a sequence.
    if topLevel.Monitors.Kind != yaml.SequenceNode {
        return fmt.Errorf("'monitors' field must be a YAML sequence")
    }

    // 3. Iterate through the sequence and decode each monitor.
    batchID := 0
    batch := make([]schema.Monitor, 0, p.config.BatchSize)

    for _, monitorNode := range topLevel.Monitors.Content {
        select {
        case <-ctx.Done():
            return ctx.Err()
        default:
            var monitor schema.Monitor
            if err := monitorNode.Decode(&monitor); err != nil {
                // Provide context for the decoding error.
                return fmt.Errorf("failed to decode monitor at line %d: %w",
monitorNode.Line, err)
            }

            // Basic validation to catch empty monitors from malformed YAML (e.g., "-").
            if monitor.Name == "" && monitor.Pulse.Type == "" {
                // This is likely an empty or malformed monitor entry, so we skip it.
                continue
            }

            batch = append(batch, monitor)

            if len(batch) >= p.config.BatchSize {
                // Send the full batch.
                batchChan <- MonitorBatch{Monitors: batch, BatchID: batchID}
                // Reset batch.
                batch = make([]schema.Monitor, 0, p.config.BatchSize)
                batchID++
            }
        }
    }
}

```

```

// 4. Send any remaining monitors in the last batch.
if len(batch) > 0 {
    batchChan <- MonitorBatch{Monitors: batch, BatchID: batchID}
}

return nil
}

/* loader/validator/errors.go */

package validator

/* loader/validator/validator.go */

package validator

//  

//import (  

//    "cptra/internal/loader/parser"  

//    parser2 "cptra/internal/loader/schema"  

//)  

//  

//type Validator interface {  

//    ValidateManifest() error  

//}  

//  

//type YamlValidator struct {  

//}  

//  

//func NewYamlValidator() *YamlValidator {  

//    return &YamlValidator{}  

//}  

//  

//func validateStringSize(s string, min int, max int) error {  

//    if len(s) < min || len(s) > max {  

//        return &parser.requiredFieldError{  

//            Field: "monitors.name",  

//            Reason: "must be between 1 and 100 characters",  

//        }
//    }
//    return nil
//}  

//  

//func (y *YamlValidator) validatePulseConfig(p parser2.PulseConfig) error {  

//  

//    switch p.(type) {  

//    case parser2.PulseHTTPConfig:  

//        if p.(parser2.PulseHTTPConfig).Url == "" {  

//            return &parser.requiredFieldError{  

//                Field: "monitors.config.url",  

//                Reason: "cannot be empty for pulse type http",
//            }
//        }
//    }
//    return nil
//}  

//  

//func (y *YamlValidator) validatePulse(p *parser2.Pulse) error {  

//    if p.Type == "" {  

//        return &parser.requiredFieldError{  

//            Field: "monitors.name",  

//            Reason: "cannot be empty",
//        }
//    }
//  

//    cfg, err := parser2.DecodePulseConfig(p)
//    if err != nil {
//        return err
//    }
//    if cfg == nil {
//        return &parser.requiredFieldError{  

//            Field: "monitors.pulse.config",  

//            Reason: "cannot be empty",
//        }
//    }
//    err = y.validatePulseConfig(cfg)
//    if err != nil {

```

```

//      return err
//    }
//    return nil
//}
//
//func (y *YamlValidator) validateMonitor(m *parser2.Monitor) error {
//    if m.Name == "" {
//        return &parser.requiredFieldError{
//            Field: "Monitor.Name",
//            Reason: "cannot be empty",
//        }
//    }
//    return nil
//}
//
//func (y *YamlValidator) ValidateManifest(m *parser2.Manifest) error {
//    for _, monitor := range m.Monitors {
//        err := y.validateMonitor(&monitor)
//        if err != nil {
//            return err
//        }
//    }
//    return nil
//}

/* loader/validator/validator_test.go */
package validator

/* main.go */
package main

import (
    "context"
    "flag"
    "fmt"
    "log"
    "os"
    "os/signal"
    "runtime"
    "runtime/pprof"
    "sync"
    "syscall"
    "time"

    "cpra/internal/controller"
)

func main() {
    // Command line flags
    var (
        configFile = flag.String("config", "", "Configuration file path")
        yamlFile   = flag.String("yaml", "internal/loader/replicated_test.yaml", "YAML file
with monitors")
        profile    = flag.Bool("profile", false, "Enable CPU profiling")
        debug      = flag.Bool("debug", false, "Enable debug logging")
    )
    flag.Parse()

    // Initialize loggers first
    controller.InitializeLoggers(*debug)

    controller.SystemLogger.Info("Starting CPRA Optimized Controller for 1M Monitors")
    controller.SystemLogger.Info("Input file: %s", *yamlFile)

    // Create optimized configuration
    config := controller.DefaultConfig()

    // Override configuration if file provided
    if *configFile != "" {
        fmt.Printf("Loading configuration from: %s\n", *configFile)
        // Configuration loading would be implemented here
    }

    // Enable profiling if requested
}

```

```

if *profile {
    fmt.Println("CPU profiling enabled")
    f, err := os.Create("cpu.pprof")
    if err != nil {
        log.Fatal("could not create CPU profile: ", err)
    }
    defer f.Close() // error handling omitted for example
    if err := pprof.StartCPUProfile(f); err != nil {
        log.Fatal("could not start CPU profile: ", err)
    }
    defer pprof.StopCPUProfile()
}

// Create the new optimized controller
oc := controller.NewOptimizedController(config)

// Setup context for graceful shutdown
ctx, cancel := context.WithCancel(context.Background())
defer cancel()

// Setup signal handler for graceful shutdown
sigChan := make(chan os.Signal, 1)
signal.Notify(sigChan, syscall.SIGINT, syscall.SIGTERM)

var shutdownInitiated bool
var shutdownMutex sync.Mutex

go func() {
    sig := <-sigChan
    shutdownMutex.Lock()
    if !shutdownInitiated {
        shutdownInitiated = true
        fmt.Printf("\nShutdown signal received (%v)...\\n", sig)
        cancel()
    }
    shutdownMutex.Unlock()
}()

// Load monitors if YAML file exists
if _, err := os.Stat(*yamlFile); err == nil {
    fmt.Printf("Loading monitors from %s...\\n", *yamlFile)
    start := time.Now()

    if err := oc.LoadMonitors(ctx, *yamlFile); err != nil {
        fmt.Printf("Error loading monitors: %v\\n", err)
        os.Exit(1)
    }

    fmt.Printf("Monitor loading completed in %v\\n", time.Since(start))
} else {
    fmt.Printf("Warning: YAML file %s not found, starting without loading monitors\\n",
*yamlFile)
}

// Start the optimized controller
if err := oc.Start(ctx); err != nil {
    fmt.Printf("Error starting controller: %v\\n", err)
    os.Exit(1)
}

// Wait for shutdown signal
<-ctx.Done()
fmt.Println("Shutting down...")

// Print memory usage
PrintMemUsage()

// Stop the controller
oc.Stop()

// Close loggers after everything is done
controller.CloseLoggers()

fmt.Println("CPRA Optimized Controller stopped")
}

// bToMb converts bytes to megabytes
func bToMb(b uint64) uint64 {
    return b / 1024 / 1024
}

```

```

// PrintMemUsage outputs the current, total, and system memory usage
func PrintMemUsage() {
    var m runtime.MemStats
    runtime.ReadMemStats(&m)
    fmt.Printf("\nMemory usage on exit:\n")
    fmt.Printf("Alloc = %v MiB", bToMb(m.Alloc))
    fmt.Printf("\tTotalAlloc = %v MiB", bToMb(m.TotalAlloc))
    fmt.Printf("\tSys = %v MiB", bToMb(m.Sys))
    fmt.Printf("\tNumGC = %v\n", m.NumGC)
}

/* queue/adaptive_queue.go */

package queue

import (
    "cpra/internal/jobs"
    "errors"
    "reflect"
    "sync/atomic"
    "time"
)

// AdaptiveQueue is a lock-free, thread-safe, fixed-size circular queue.
// It is designed for high-throughput scenarios with multiple producers and consumers.
// It implements the Queue interface.
type AdaptiveQueue struct {
    buffer      []jobs.Job
    capacity    uint64
    head        uint64
    tail        uint64
    closed      int32

    enqueueuedCount     int64
    dequeuedCount       int64
    totalQueueWaitNanos int64
    maxQueueWaitNanos   int64
    startUnixNano       int64
    lastEnqueueUnixNano int64
    lastDequeueUnixNano int64
}

// NewAdaptiveQueue creates a new AdaptiveQueue with the given capacity.
// Capacity must be a power of 2 for efficient bitwise operations.
func NewAdaptiveQueue(capacity uint64) (*AdaptiveQueue, error) {
    if (capacity & (capacity - 1)) != 0 {
        return nil, errors.New("capacity must be a power of 2")
    }
    queue := &AdaptiveQueue{
        buffer:      make([]jobs.Job, capacity),
        capacity:    capacity,
        startUnixNano: time.Now().UnixNano(),
    }
    return queue, nil
}

// Enqueue adds a single job to the queue.
func (q *AdaptiveQueue) Enqueue(job jobs.Job) error {
    if atomic.LoadInt32(&q.closed) == 1 {
        return ErrQueueClosed
    }

    now := time.Now()
    for {
        head := atomic.LoadUint64(&q.head)
        tail := atomic.LoadUint64(&q.tail)

        if tail-head >= q.capacity {
            return ErrQueueFull // Queue is full
        }

        // Attempt to claim the next spot
        if atomic.CompareAndSwapUint64(&q.tail, tail, tail+1) {
            if !isNilJob(job) {
                job.SetEnqueueTime(now)
            }
            q.buffer[tail&(q.capacity-1)] = job
            atomic.AddInt64(&q.enqueueuedCount, 1)
        }
    }
}

```

```

        atomic.StoreInt64(&q.lastEnqueueUnixNano, now.UnixNano())
        return nil
    }
}

// EnqueueBatch adds a batch of jobs to the queue using a highly concurrent, lock-free algorithm.
func (q *AdaptiveQueue) EnqueueBatch(jobsInterface []interface{}) error {
    if len(jobsInterface) == 0 {
        return nil
    }

    // Convert interface{} slice to jobs.Job slice
    convertedJobs := make([]jobs.Job, len(jobsInterface))
    for i, job := range jobsInterface {
        if j, ok := job.(jobs.Job); ok {
            convertedJobs[i] = j
        } else {
            return errors.New("invalid job type in batch")
        }
    }
    if atomic.LoadInt32(&q.closed) == 1 {
        return ErrQueueClosed
    }
    n := uint64(len(convertedJobs))

    now := time.Now()
    for {
        head := atomic.LoadUint64(&q.head)
        tail := atomic.LoadUint64(&q.tail)

        if tail-head+n > q.capacity {
            return ErrQueueFull
        }

        // Atomically claim a slot for the entire batch
        if atomic.CompareAndSwapUint64(&q.tail, tail, tail+n) {
            // Once the slot is claimed, we can write the batch without further atomics
            for i := uint64(0); i < n; i++ {
                job := convertedJobs[i]
                if !isNilJob(job) {
                    job.SetEnqueueTime(now)
                }
                q.buffer[(tail+i)&(q.capacity-1)] = job
            }
            atomic.AddInt64(&q.enqueuedCount, int64(n))
            atomic.StoreInt64(&q.lastEnqueueUnixNano, now.UnixNano())
            return nil
        }
        // If CAS fails, another producer got there first. Loop and try again.
    }
}

// DequeueBatch removes and returns a batch of jobs from the queue.
func (q *AdaptiveQueue) DequeueBatch(maxSize int) ([]jobs.Job, error) {
    if atomic.LoadInt32(&q.closed) == 1 && q.IsEmpty() {
        return nil, ErrQueueClosed
    }

    for {
        head := atomic.LoadUint64(&q.head)
        tail := atomic.LoadUint64(&q.tail)

        if head >= tail {
            return nil, nil // Queue is empty
        }

        n := tail - head
        if n > uint64(maxSize) {
            n = uint64(maxSize)
        }

        // Atomically claim the batch for dequeuing
        if atomic.CompareAndSwapUint64(&q.head, head, head+n) {
            batch := make([]jobs.Job, n)
            now := time.Now()
            for i := uint64(0); i < n; i++ {
                batch[i] = q.buffer[(head+i)&(q.capacity-1)]
                // Nil out the buffer slot to help the GC
                q.buffer[(head+i)&(q.capacity-1)] = nil
            }
            atomic.AddInt64(&q.dequeuedCount, int64(n))
            atomic.StoreInt64(&q.lastDequeueUnixNano, now.UnixNano())
            return batch, nil
        }
    }
}

```

```

        if !isNilJob(batch[i]) {
            enqueueTime := batch[i].GetEnqueueTime()
            if !enqueueTime.IsZero() {
                wait := now.Sub(enqueueTime)
                atomic.AddInt64(&q.totalQueueWaitNanos, int64(wait))
                for {
                    currentMax :=
                        if waitNs := int64(wait); waitNs <= currentMax
                            break
                }
                if
atomic.CompareAndSwapInt64(&q.maxQueueWaitNanos, currentMax, int64(wait)) {
                                break
                            }
                        }
                    }
                }
            }
        atomic.AddInt64(&q.dequeuedCount, int64(n))
        atomic.StoreInt64(&q.lastDequeueUnixNano, now.UnixNano())
        return batch, nil
    }
    // If CAS fails, another consumer got there first. Loop and try again.
}
}

// Dequeue removes and returns a single job from the queue.
func (q *AdaptiveQueue) Dequeue() (jobs.Job, error) {
    if atomic.LoadInt32(&q.closed) == 1 && q.IsEmpty() {
        return nil, ErrQueueClosed
    }

    for {
        head := atomic.LoadUint64(&q.head)
        tail := atomic.LoadUint64(&q.tail)

        if head >= tail {
            return nil, nil // Queue is empty
        }

        job := q.buffer[head&(q.capacity-1)]

        // Attempt to move the head pointer
        if atomic.CompareAndSwapUint64(&q.head, head, head+1) {
            now := time.Now()
            if !isNilJob(job) {
                enqueueTime := job.GetEnqueueTime()
                if !enqueueTime.IsZero() {
                    wait := now.Sub(enqueueTime)
                    atomic.AddInt64(&q.totalQueueWaitNanos, int64(wait))
                    for {
                        currentMax := atomic.LoadInt64(&q.maxQueueWaitNanos)
                        if waitNs := int64(wait); waitNs <= currentMax {
                            break
                        }
                        if atomic.CompareAndSwapInt64(&q.maxQueueWaitNanos,
currentMax, int64(wait)) {
                                break
                            }
                        }
                    }
                }
            }
            atomic.AddInt64(&q.dequeuedCount, 1)
            atomic.StoreInt64(&q.lastDequeueUnixNano, now.UnixNano())
            return job, nil
        }
    }
}

// IsEmpty checks if the queue is empty.
func (q *AdaptiveQueue) IsEmpty() bool {
    return atomic.LoadUint64(&q.head) == atomic.LoadUint64(&q.tail)
}

// Close marks the queue as closed.
func (q *AdaptiveQueue) Close() {
    atomic.StoreInt32(&q.closed, 1)
}

```

```

}
func isNilJob(job jobs.Job) bool {
    if job == nil {
        return true
    }
    val := reflect.ValueOf(job)
    switch val.Kind() {
    case reflect.Ptr, reflect.Interface, reflect.Slice, reflect.Map, reflect.Func, reflect.Chan:
        return val.IsNil()
    default:
        return false
    }
}

// Stats returns the current statistics for the queue.
// Note: This is a simplified version. A full implementation would track more metrics.
func (q *AdaptiveQueue) Stats() Stats {
    head := atomic.LoadUint64(&q.head)
    tail := atomic.LoadUint64(&q.tail)
    depth := tail - head
    enq := atomic.LoadInt64(&q.enqueuedCount)
    deq := atomic.LoadInt64(&q.dequeuedCount)
    elapsed := time.Since(time.Unix(0, atomic.LoadInt64(&q.startUnixNano)))
    if elapsed <= 0 {
        elapsed = time.Millisecond
    }
    avgWaitNs := int64(0)
    if deq > 0 {
        avgWaitNs = atomic.LoadInt64(&q.totalQueueWaitNanos) / deq
    }
    stats := Stats{
        QueueDepth:    int(depth),
        Capacity:     int(q.capacity),
        Enqueued:      enq,
        Dequeued:      deq,
        Dropped:       0,
        MaxQueueTime:  time.Duration	atomic.LoadInt64(&q.maxQueueWaitNanos)),
        AvgQueueTime:  time.Duration	avgWaitNs,
        EnqueueRate:   float64(enq) / elapsed.Seconds(),
        DequeueRate:   float64(deq) / elapsed.Seconds(),
        LastEnqueue:   time.Unix(0, atomic.LoadInt64(&q.lastEnqueueUnixNano)),
        LastDequeue:   time.Unix(0, atomic.LoadInt64(&q.lastDequeueUnixNano)),
        SampleWindow:  elapsed,
    }
    return stats
}

```

/* queue/bounded_queue.go */

```

package queue

import (
    "errors"
    "sync"
    "sync/atomic"
    "time"

    "cpra/internal/jobs"
)

var (
    ErrQueueFull    = errors.New("queue is full")
    ErrQueueClosed = errors.New("queue is closed")
)

```

// BoundedQueue implements a high-performance bounded queue with batching.
// It now implements the queue.Queue interface.

```

type BoundedQueue struct {
    batches      chan []jobs.Job
    maxSize      int32
    maxBatch     int32
    closed       int32
    enqueueued  int64
    dequeueued  int64
    dropped      int64
    mu           sync.RWMutex
    batchTimeout time.Duration
}

```

```

// BoundedQueueConfig holds queue configuration.
type BoundedQueueConfig struct {
    MaxSize      int
    MaxBatch     int
    BatchTimeout time.Duration
}

// NewBoundedQueue creates a new bounded queue.
func NewBoundedQueue(config BoundedQueueConfig) *BoundedQueue {
    return &BoundedQueue{
        batches:      make(chan []jobs.Job, config.MaxSize),
        maxSize:     int32(config.MaxSize),
        maxBatch:    int32(config.MaxBatch),
        batchTimeout: config.BatchTimeout,
    }
}

// Enqueue adds a single job to the queue.
func (q *BoundedQueue) Enqueue(job jobs.Job) error {
    return q.EnqueueBatch([]jobs.Job{job})
}

// EnqueueBatch adds a batch of jobs to the queue.
func (q *BoundedQueue) EnqueueBatch(jobs []jobs.Job) error {
    if atomic.LoadInt32(&q.closed) == 1 {
        return ErrQueueClosed
    }

    if len(jobs) == 0 {
        return nil
    }

    enqueueTime := time.Now()
    for _, job := range jobs {
        job.SetEnqueueTime(enqueueTime)
    }

    batch := jobs
    if len(batch) > int(q.maxBatch) {
        batch = batch[:q.maxBatch]
    }

    select {
    case q.batches <- batch:
        atomic.AddInt64(&q.enqueued, int64(len(batch)))
        return nil
    default:
        atomic.AddInt64(&q.dropped, int64(len(batch)))
        return ErrQueueFull
    }
}

// Dequeue removes and returns a single job from the queue.
func (q *BoundedQueue) Dequeue() (jobs.Job, error) {
    // This is inefficient for a bounded queue, but it satisfies the interface.
    // The adaptive queue will have a proper single-item dequeue.
    select {
    case batch, ok := <-q.batches:
        if !ok {
            return nil, ErrQueueClosed
        }
        atomic.AddInt64(&q.dequeued, int64(len(batch)))
        if len(batch) > 1 {
            // Re-enqueue the rest of the batch. This is very inefficient.
            go func() { q.batches <- batch[1:] }()
        }
        return batch[0], nil
    case <-time.After(10 * time.Millisecond): // Non-blocking with a small timeout
        return nil, nil // Queue is empty
    }
}

// DequeueBatch removes a batch of jobs from the queue.
func (q *BoundedQueue) DequeueBatch(maxSize int) ([]jobs.Job, error) {
    select {
    case batch, ok := <-q.batches:
        if !ok {
            return nil, ErrQueueClosed
        }

```

```

        atomic.AddInt64(&q.dequeued, int64(len(batch)))
        if len(batch) > maxSize {
            // This is inefficient, but necessary to respect maxSize.
            go func() { q.batches <- batch[maxSize:] }()
            return batch[:maxSize], nil
        }
        return batch, nil
    default:
        // Non-blocking, return empty if no batch is immediately available.
        return nil, nil
    }
}

// Close closes the queue.
func (q *BoundedQueue) Close() {
    if atomic.CompareAndSwapInt32(&q.closed, 0, 1) {
        close(q.batches)
    }
}

// Stats returns current queue statistics.
func (q *BoundedQueue) Stats() Stats {
    return Stats{
        Enqueued:    atomic.LoadInt64(&q.enqueued),
        Dequeued:    atomic.LoadInt64(&q.dequeued),
        Dropped:     atomic.LoadInt64(&q.dropped),
        QueueDepth:  len(q.batches),
        Capacity:    int(q.maxSize),
    }
}

/* queue/dynamic_worker_pool.go */

package queue

import (
    "context"
    "cpra/internal/jobs"
    "log"
    "math"
    "sync"
    "sync/atomic"
    "time"

    "github.com/panjf2000/ants/v2"
)

// ResultRouter handles routing of job results to type-specific channels.
// This enables decoupling result processing from the main worker pool.
type ResultRouter struct {
    PulseResultChan      chan []jobs.Result
    InterventionResultChan chan []jobs.Result
    CodeResultChan       chan []jobs.Result

    config WorkerPoolConfig
    logger *log.Logger
}

// WorkerPoolStats exposes runtime metrics for the dynamic worker pool.
type WorkerPoolStats struct {
    MinWorkers      int
    MaxWorkers      int
    CurrentCapacity int
    RunningWorkers   int
    WaitingTasks     int
    TargetWorkers    int
    TasksSubmitted   int64
    TasksCompleted   int64
    ScalingEvents    int64
    LastScaleTime    time.Time
    PendingResults   int
}

// NewResultRouter creates a new result router with buffered channels.
func NewResultRouter(config WorkerPoolConfig, logger *log.Logger) *ResultRouter {
    bufferSize := config.MaxWorkers // Buffer size based on max workers
    return &ResultRouter{
        PulseResultChan:      make(chan []jobs.Result, bufferSize),
        InterventionResultChan: make(chan []jobs.Result, bufferSize),

```

```

        CodeResultChan:      make(chan []jobs.Result, bufferSize),
        config:            config,
        logger:           logger,
    }
}

// RouteResults takes a batch of mixed results and routes them to appropriate channels.
func (r *ResultRouter) RouteResults(results []jobs.Result) {
    if len(results) == 0 {
        return
    }

    // Group results by type
    pulseResults := make([][]jobs.Result, 0, len(results))
    interventionResults := make([][]jobs.Result, 0, len(results))
    codeResults := make([][]jobs.Result, 0, len(results))

    for _, result := range results {
        switch result.Payload["type"] {
        case "pulse":
            pulseResults = append(pulseResults, result)
        case "intervention":
            interventionResults = append(interventionResults, result)
        case "code":
            codeResults = append(codeResults, result)
        default:
            r.logger.Printf("Unknown job type in result: %v", result.Payload["type"])
        }
    }

    // Send to appropriate channels with backpressure logging
    if len(pulseResults) > 0 {
        r.sendWithBackpressure(r.PulseResultChan, pulseResults, "pulse")
    }
    if len(interventionResults) > 0 {
        r.sendWithBackpressure(r.InterventionResultChan, interventionResults, "intervention")
    }
    if len(codeResults) > 0 {
        r.sendWithBackpressure(r.CodeResultChan, codeResults, "code")
    }
}

func (r *ResultRouter) sendWithBackpressure(ch chan []jobs.Result, batch []jobs.Result, label string) {
    backoff := r.config.ResultBatchTimeout
    if backoff <= 0 {
        backoff = 50 * time.Millisecond
    }
    ticker := time.NewTicker(backoff)
    defer ticker.Stop()

    for {
        select {
        case ch <- batch:
            return
        case <-ticker.C:
            r.logger.Printf("Backpressure: %s results stalled (%d jobs waiting)", label,
len(batch))
        }
    }
}

// Close closes all result channels.
func (r *ResultRouter) Close() {
    close(r.PulseResultChan)
    close(r.InterventionResultChan)
    close(r.CodeResultChan)
}

// DynamicWorkerPool manages a pool of workers that execute jobs from a queue.
// It can dynamically adjust the number of workers based on load.
type DynamicWorkerPool struct {
    queue      Queue
    antsPool   *ants.PoolWithFunc
    logger     *log.Logger
    config     WorkerPoolConfig
    resultChan chan jobs.Result
    router     *ResultRouter
    ctx       context.Context
}

```

```

cancel context.CancelFunc
wg     sync.WaitGroup

tasksSubmitted int64
tasksCompleted int64
scalingEvents  int64
lastTarget    int64
lastScaleTime int64
stopping      int32
}

// WorkerPoolConfig holds configuration for the DynamicWorkerPool.
type WorkerPoolConfig struct {
    MinWorkers      int
    MaxWorkers      int
    AdjustmentInterval time.Duration
    ResultBatchSize  int
    ResultBatchTimeout time.Duration
    TargetQueueLatency time.Duration
}

// DefaultWorkerPoolConfig returns a default configuration for the worker pool.
func DefaultWorkerPoolConfig() WorkerPoolConfig {
    return WorkerPoolConfig{
        MinWorkers:      10,
        MaxWorkers:     10000,
        AdjustmentInterval: 5 * time.Second,
        ResultBatchSize: 1000,
        ResultBatchTimeout: 10 * time.Millisecond,
        TargetQueueLatency: 100 * time.Millisecond,
    }
}

// NewDynamicWorkerPool creates a new dynamic worker pool.
func NewDynamicWorkerPool(q Queue, config WorkerPoolConfig, logger *log.Logger) (*DynamicWorkerPool, error) {
    if config.MinWorkers <= 0 {
        config.MinWorkers = 1
    }
    if config.MaxWorkers < config.MinWorkers {
        config.MaxWorkers = config.MinWorkers
    }
    if config.ResultBatchSize <= 0 {
        config.ResultBatchSize = config.MaxWorkers
    }
    if config.TargetQueueLatency <= 0 {
        config.TargetQueueLatency = 100 * time.Millisecond
    }

    ctx, cancel := context.WithCancel(context.Background())

    pool := &DynamicWorkerPool{
        queue:      q,
        logger:     logger,
        config:     config,
        resultChan: make(chan jobs.Result, config.MaxWorkers),
        router:     NewResultRouter(config, logger),
        ctx:        ctx,
        cancel:    cancel,
    }

    workerFunc := func(job interface{}) {
        j, ok := job.(jobs.Job)
        if !ok {
            pool.logger.Printf("Error: Invalid job type in worker pool: %T", job)
            return
        }
        result := j.Execute()
        if atomic.LoadInt32(&pool.stopping) == 1 {
            return
        }
        select {
        case pool.resultChan <- result:
        case <-pool.ctx.Done():
        }
    }

    antsPool, err := ants.NewPoolWithFunc(config.MaxWorkers, workerFunc,
ants.WithPanicHandler(func(err interface{}) {
    pool.logger.Printf("Worker panic: %v", err)
})
}

```

```

    }})
    if err != nil {
        return nil, err
    }
    pool.antsPool = antsPool
    pool.antsPool.Tune(config.MinWorkers)
    atomic.StoreInt64(&pool.lastTarget, int64(config.MinWorkers))
    atomic.StoreInt64(&pool.lastScaleTime, time.Now().UnixNano())

    return pool, nil
}

// Start begins the worker pool's operations.
func (p *DynamicWorkerPool) Start() {
    routineCount := 2
    if p.config.AdjustmentInterval > 0 {
        routineCount++
    }
    p.wg.Add(routineCount)
    go p.dispatcher()
    go p.resultProcessor()
    if p.config.AdjustmentInterval > 0 {
        go p.autoScale()
    }
    p.logger.Println("DynamicWorkerPool started")
}

// GetRouter returns the result router for accessing type-specific result channels.
func (p *DynamicWorkerPool) GetRouter() *ResultRouter {
    return p.router
}

// DrainAndStop waits for outstanding tasks to finish before stopping the worker pool.
func (p *DynamicWorkerPool) DrainAndStop() {
    if !atomic.CompareAndSwapInt32(&p.stopping, 0, 1) {
        return
    }
    p.logger.Println("Draining DynamicWorkerPool...")
    p.cancel()
    done := make(chan struct{})
    go func() {
        p.wg.Wait()
        close(done)
    }()
    select {
    case <-done:
    case <-time.After(p.config.TargetQueueLatency * 5):
        p.logger.Println("Draining timed out, continuing shutdown")
    }
    remaining := len(p.resultChan)
    if remaining > 0 {
        p.logger.Printf("Flushing %d queued results before close", remaining)
    }
    close(p.resultChan)
    p.router.Close()
    p.antsPool.Release()
    p.logger.Println("DynamicWorkerPool stopped")
}

// dispatcher fetches batches of jobs from the queue and submits them to the ants pool.
func (p *DynamicWorkerPool) dispatcher() {
    defer p.wg.Done()
    for {
        select {
        case <-p.ctx.Done():
            return
        default:
            batchTarget := p.antsPool.Cap()
            if batchTarget <= 0 {
                batchTarget = p.config.MinWorkers
            }
            if batchTarget > p.config.ResultBatchSize {
                batchTarget = p.config.ResultBatchSize
            }
            if batchTarget <= 0 {
                batchTarget = 1
            }

            jobs, err := p.queue.DequeueBatch(batchTarget)
            if err != nil {

```

```

        if err != ErrQueueClosed {
            p.logger.Printf("Error dequeuing job batch: %v", err)
        }
        time.Sleep(100 * time.Millisecond) // Wait a bit if there's an error
        continue
    }
    if len(jobs) == 0 {
        time.Sleep(10 * time.Millisecond) // Wait if the queue is empty
        continue
    }

    atomic.AddInt64(&p.tasksSubmitted, int64(len(jobs)))

    for _, job := range jobs {
        if err := p.antsPool.Invoke(job); err != nil {
            p.logger.Printf("Error invoking job: %v", err)
        }
    }
}

// resultProcessor collects individual results and routes them through the router in batches.
func (p *DynamicWorkerPool) resultProcessor() {
    defer p.wg.Done()

    batch := make([]jobs.Result, 0, p.config.ResultBatchSize)
    ticker := time.NewTicker(p.config.ResultBatchTimeout)
    defer ticker.Stop()

    for {
        select {
        case <-p.ctx.Done():
            // Route any remaining results before shutting down
            if len(batch) > 0 {
                p.router.RouteResults(batch)
            }
            return
        case result, ok := <-p.resultChan:
            if !ok { // resultChan was closed
                if len(batch) > 0 {
                    p.router.RouteResults(batch)
                }
                return
            }
            atomic.AddInt64(&p.tasksCompleted, 1)
            batch = append(batch, result)
            if len(batch) >= p.config.ResultBatchSize {
                p.router.RouteResults(batch)
                batch = make([]jobs.Result, 0, p.config.ResultBatchSize)
                // Reset the ticker to prevent immediate firing
                ticker.Reset(p.config.ResultBatchTimeout)
            }
        case <-ticker.C:
            // Route partial batches on timeout
            if len(batch) > 0 {
                p.router.RouteResults(batch)
                batch = make([]jobs.Result, 0, p.config.ResultBatchSize)
            }
        }
    }
}

// autoScale periodically tunes the ants pool capacity based on queue depth.
func (p *DynamicWorkerPool) autoScale() {
    defer p.wg.Done()

    ticker := time.NewTicker(p.config.AdjustmentInterval)
    defer ticker.Stop()

    for {
        select {
        case <-p.ctx.Done():
            return
        case <-ticker.C:
            stats := p.queue.Stats()
            desired := p.desiredCapacity(stats)
            current := p.antsPool.Cap()
            if desired != current {
                p.antsPool.Tune(desired)
            }
        }
    }
}

```

```

desired, stats.QueueDepth)
    p.logger.Printf("Tuned worker pool capacity to %d (queue depth=%d)",
        atomic.StoreInt64(&p.lastTarget, int64(desired))
        atomic.StoreInt64(&p.lastScaleTime, time.Now().UnixNano())
        atomic.AddInt64(&p.scalingEvents, 1)
    }
}
}

func (p *DynamicWorkerPool) desiredCapacity(stats QueueStats) int {
    current := p.antsPool.Cap()
    if current <= 0 {
        current = p.config.MinWorkers
    }

    minWorkers := p.config.MinWorkers
    maxWorkers := p.config.MaxWorkers
    if maxWorkers < minWorkers {
        maxWorkers = minWorkers
    }

    enqueueRate := stats.EnqueueRate
    if enqueueRate <= 0 {
        enqueueRate = stats.DequeueRate
    }
    targetLatency := p.config.TargetQueueLatency
    if targetLatency <= 0 {
        targetLatency = 100 * time.Millisecond
    }

    desired := current

    // Estimate per-worker throughput
    perWorker := 0.0
    if current > 0 && stats.DequeueRate > 0 {
        perWorker = stats.DequeueRate / float64(current)
    }
    if perWorker > 0 && enqueueRate > 0 {
        desired = int(math.Ceil(enqueueRate / perWorker))
    }

    // Enforce latency budget using Little's Law ( $L = \lambda W$ )
    if enqueueRate > 0 {
        targetDepth := enqueueRate * targetLatency.Seconds()
        // Always allow at least minWorkers entities worth of backlog
        if targetDepth < float64(minWorkers) {
            targetDepth = float64(minWorkers)
        }
        depth := float64(stats.QueueDepth)
        if depth > targetDepth && targetDepth > 0 {
            scale := depth / targetDepth
            desired = int(math.Ceil(float64(desired) * scale))
        } else if depth < targetDepth/2 && desired > minWorkers {
            desired = int(math.Max(float64(minWorkers), math.Ceil(float64(desired)*0.8)))
        }
    }

    if desired < minWorkers {
        desired = minWorkers
    }
    if desired > maxWorkers {
        desired = maxWorkers
    }
    return desired
}

// Stats returns runtime statistics for the worker pool.
func (p *DynamicWorkerPool) Stats() WorkerPoolStats {
    return WorkerPoolStats{
        MinWorkers:      p.config.MinWorkers,
        MaxWorkers:      p.config.MaxWorkers,
        CurrentCapacity: p.antsPool.Cap(),
        RunningWorkers: p.antsPool.Running(),
        WaitingTasks:   p.antsPool.Waiting(),
        TargetWorkers:   int	atomic.LoadInt64(&p.lastTarget),
        TasksSubmitted: atomic.LoadInt64(&p.tasksSubmitted),
        TasksCompleted: atomic.LoadInt64(&p.tasksCompleted),
        ScalingEvents:  atomic.LoadInt64(&p.scalingEvents),
        LastScaleTime:   time.Unix(0, atomic.LoadInt64(&p.lastScaleTime)),
    }
}

```

```

        PendingResults: len(p.resultChan),
    }

/* queue/queue.go */

package queue

import (
    "cpra/internal/jobs"
    "time"
)

// Queue defines the interface for a generic, thread-safe queue system.
// This allows the controller and systems to be decoupled from a specific queue implementation.
type Queue interface {
    // Enqueue adds a single job to the queue.
    Enqueue(job jobs.Job) error

    // EnqueueBatch adds a slice of jobs to the queue.
    EnqueueBatch(jobs []interface{}) error

    // Dequeue removes and returns a single job from the queue.
    Dequeue() (jobs.Job, error)

    // DequeueBatch removes and returns a batch of jobs from the queue.
    DequeueBatch(maxSize int) ([]jobs.Job, error)

    // Close shuts down the queue and prevents new jobs from being enqueued.
    Close()

    // Stats returns statistics about the queue's performance.
    Stats() Stats
}

// Stats holds performance metrics for a queue.
type Stats struct {
    QueueDepth      int
    Capacity        int
    Enqueued        int64
    Dequeued        int64
    Dropped          int64
    MaxQueueTime   time.Duration
    AvgQueueTime   time.Duration
    MaxJobLatency  time.Duration
    AvgJobLatency  time.Duration
    EnqueueRate     float64 // jobs per second since queue creation
    DequeueRate     float64 // jobs per second since queue creation
    LastEnqueue     time.Time
    LastDequeue     time.Time
    SampleWindow    time.Duration // elapsed time used for rate calculations
}

```